

CS 330: Network Applications & Protocols

Transport Layer

Galin Zhelezov
Department of Physical Sciences
York College of Pennsylvania



Overview of Transport Layer

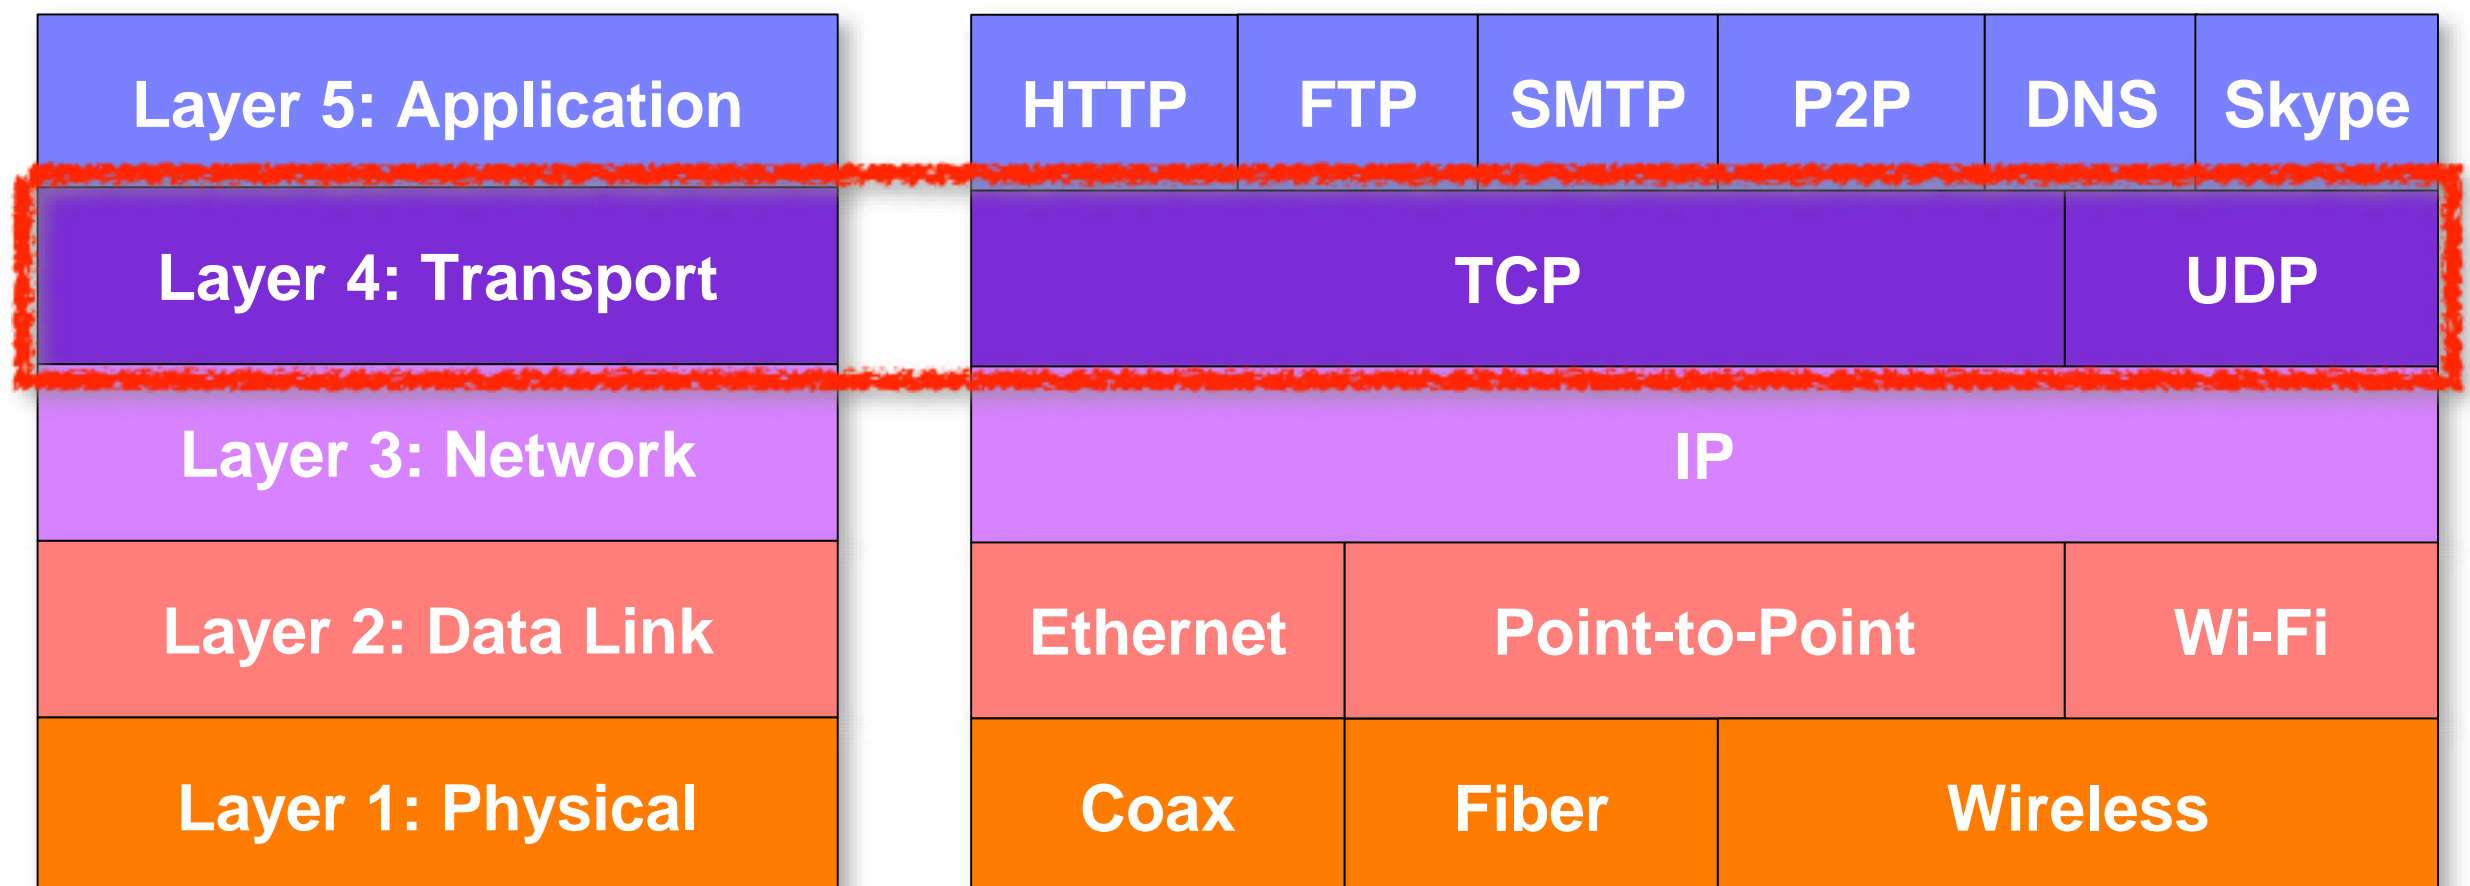
- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
- **Principles of Congestion Control**
- **TCP Congestion Control**

Overview of Transport Layer


- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
- **Principles of Congestion Control**
- **TCP Congestion Control**

Protocol Layers

- **Top-Down Approach**



Transport Layer

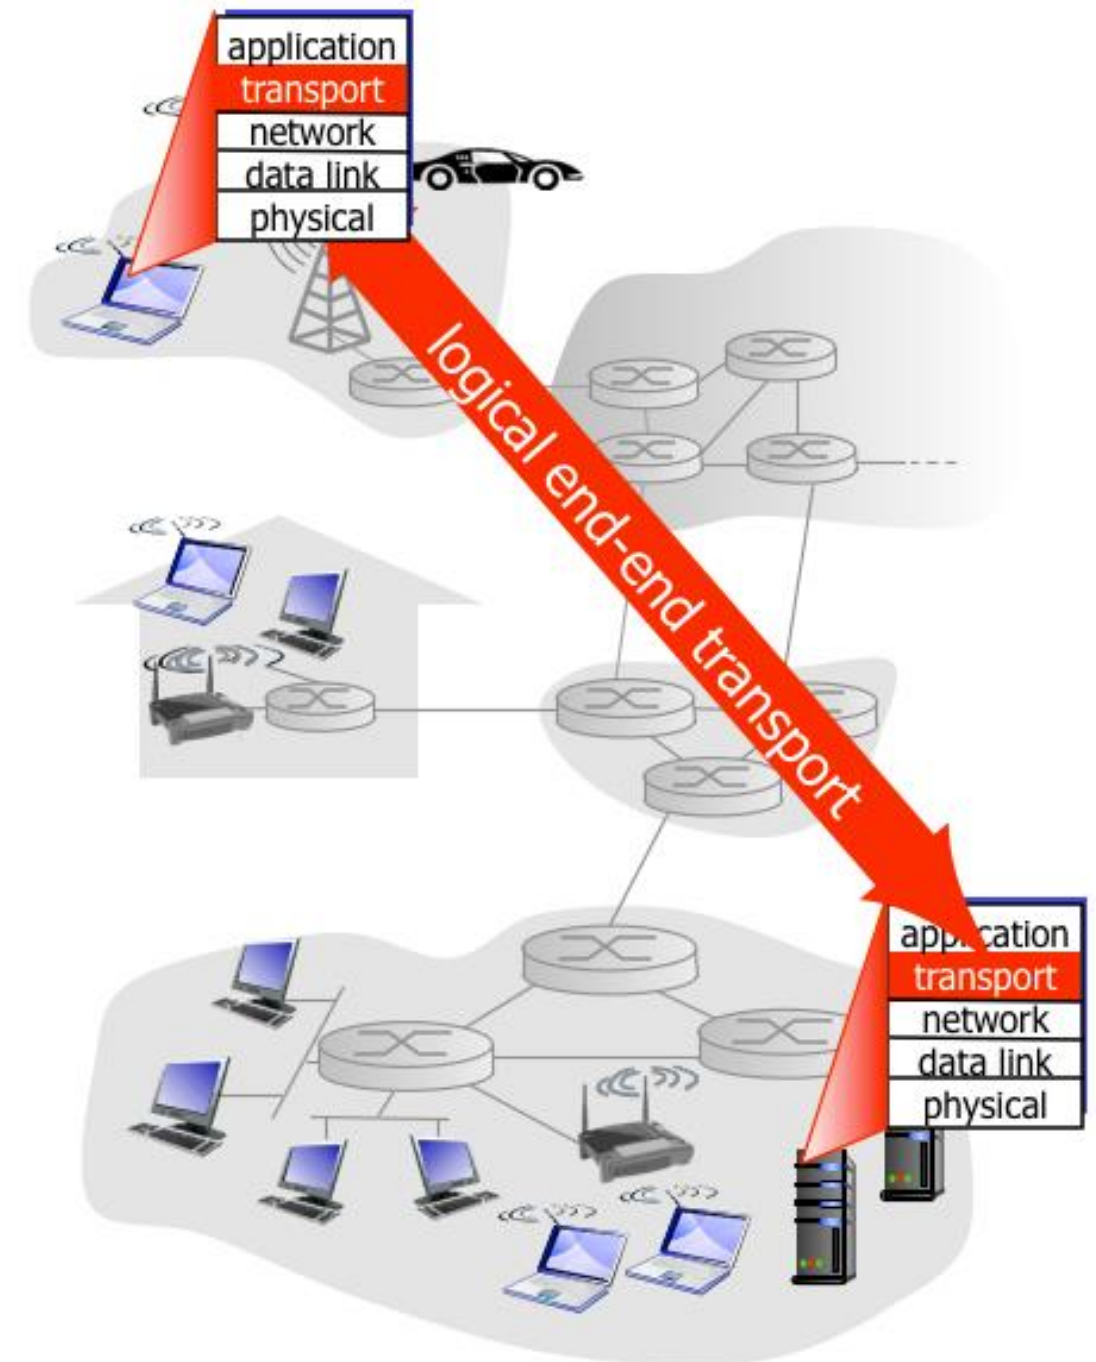


- **Transport layer provides End-to-End Services**

- Required at source and destination
- Not required at intermediate hops

Transport Services and Protocols

- **Provide logical communication between application processes running on different hosts**
- **Transport protocols run in end systems**
 - **Sending side**: breaks application messages into segments, passes segments down to network layer
 - **Receiving side**: reassembles segments into messages, passes to application layer
- **More than one transport protocol available to apps**
 - Internet: TCP and UDP

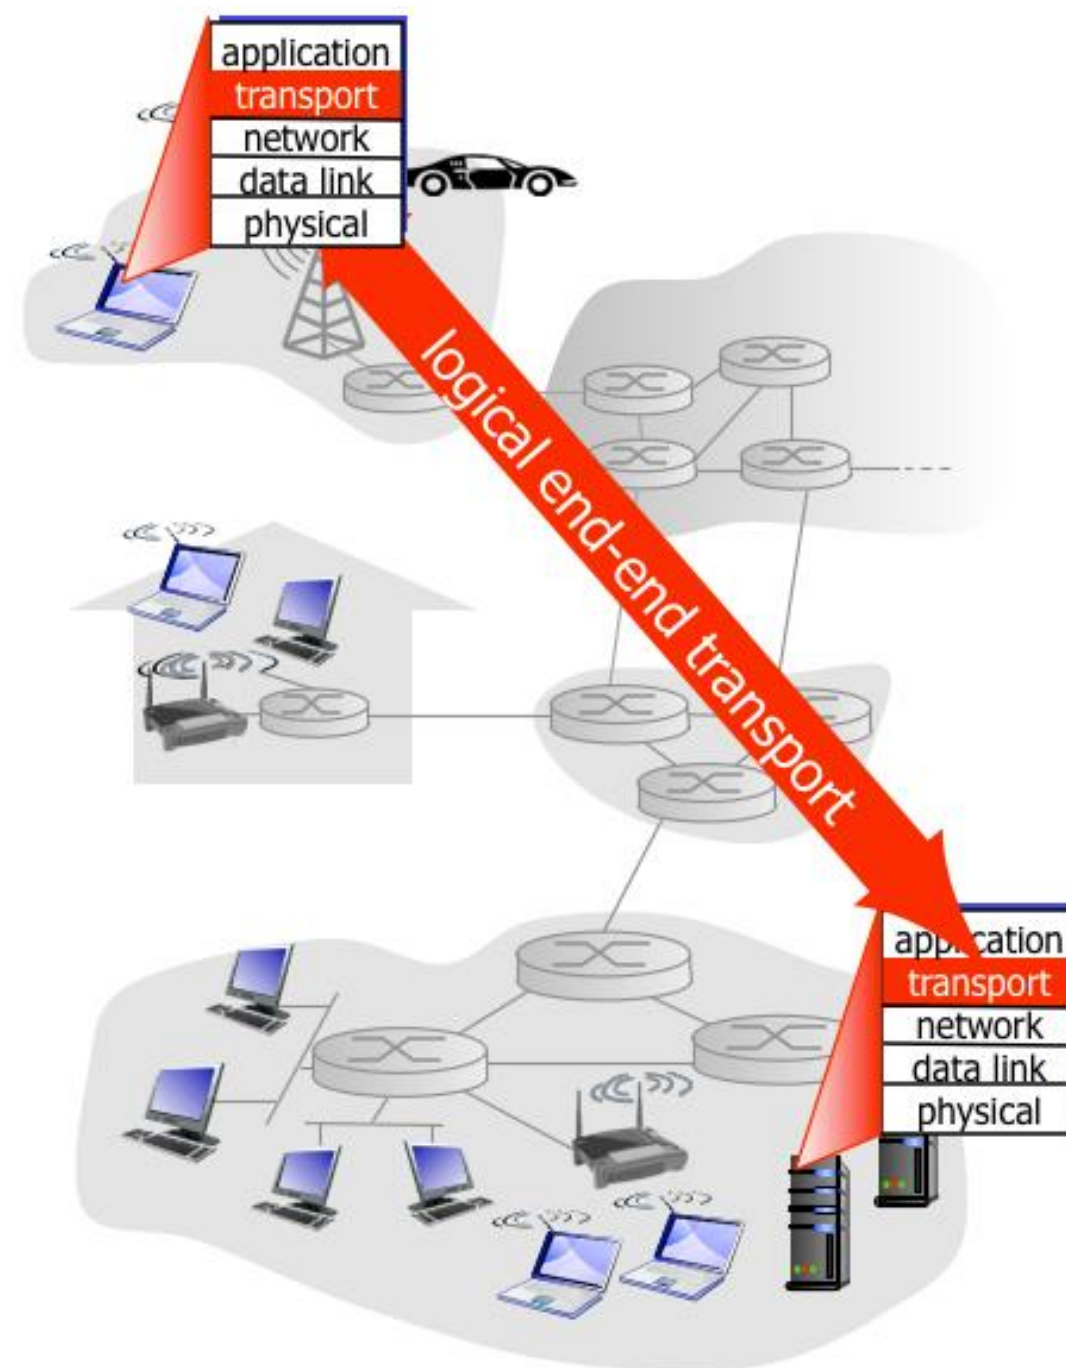


Transport vs. Network layer

- **Network layer: logical communication between hosts**
- **Transport layer: logical communication between processes**
 - Relies on, enhances, network layer services

Internet Transport-layer Protocols

- **Reliable, in-order delivery: TCP**
 - Congestion control
 - Flow control
 - Connection setup
- **Unreliable, unordered delivery: UDP**
 - No-frills extension of “best-effort” IP
- **Services not available:**
 - Delay guarantees
 - Bandwidth guarantees



Transport Layer Functions

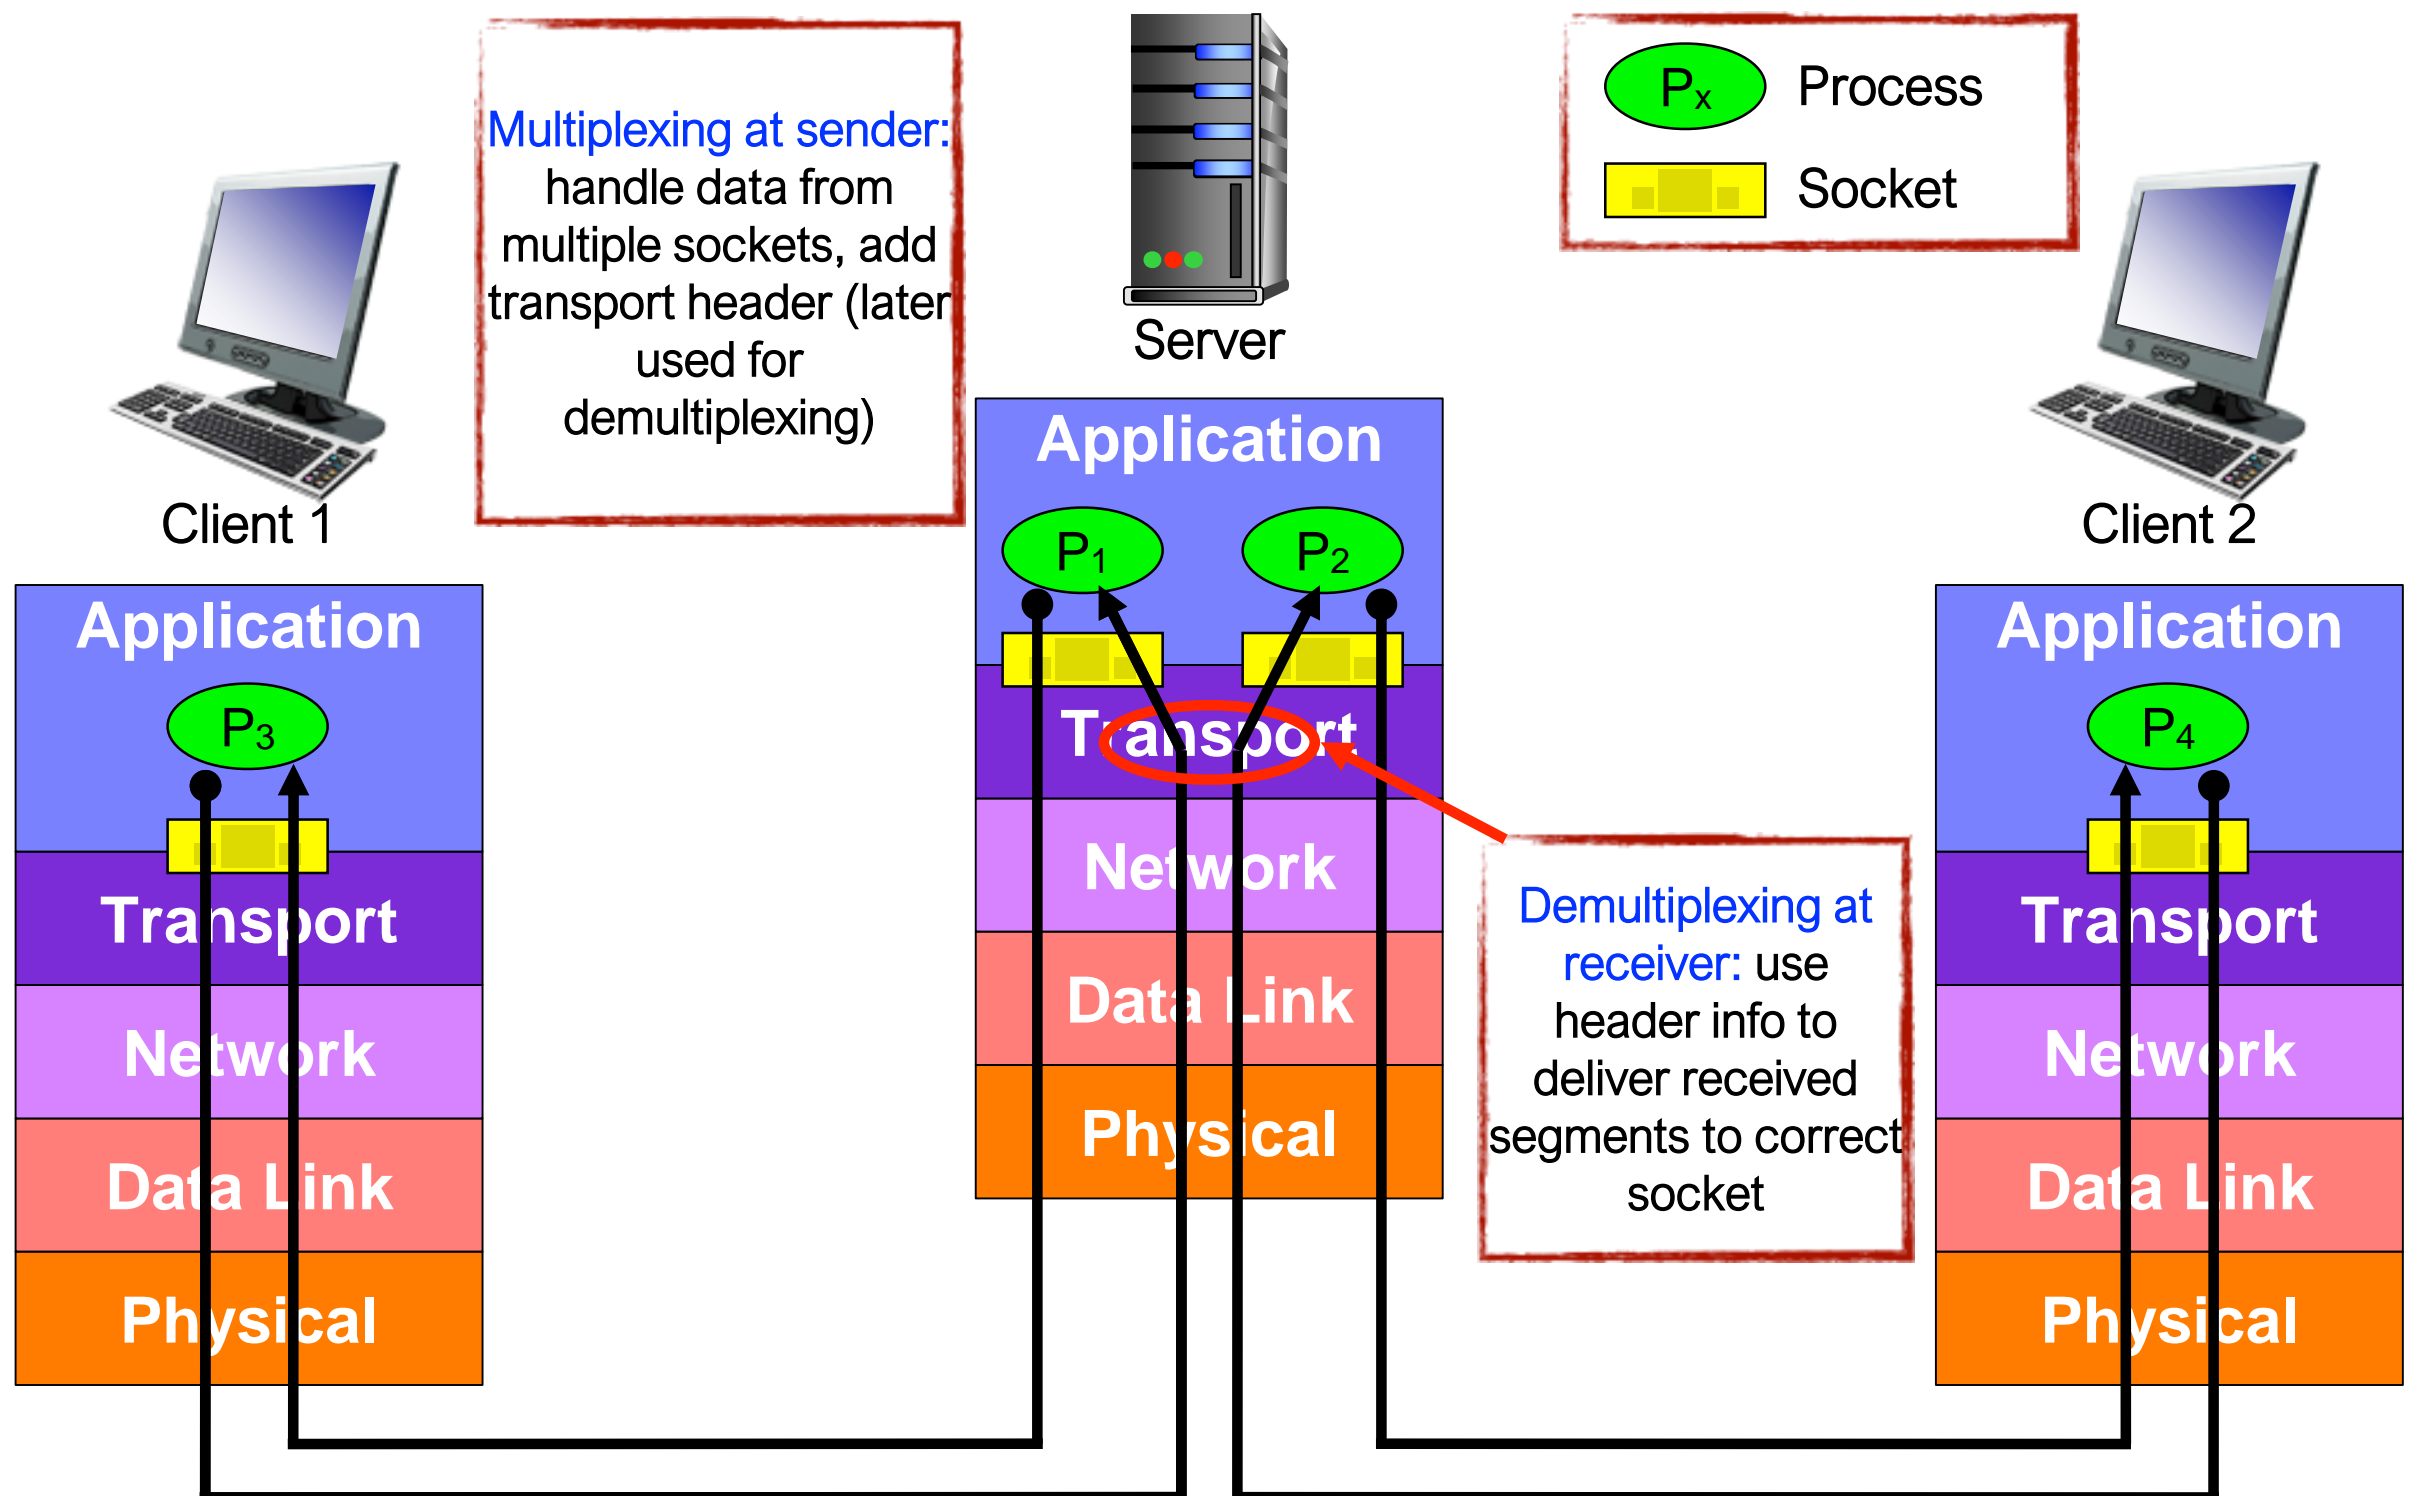
- **Multiplexing and demultiplexing** among applications and processes at end systems
- **Error detection** of bit errors
- **Loss detection** - lost packets due to buffer overflow at intermediate systems
- **Error/loss Recovery** - retransmissions
- **Flow Control** - ensures receiver has buffer capacity to receive message
- **Congestion Control** - ensure the network has the capacity to transmit data

Not all transport protocols provide all of the above functionality

Overview of Transport Layer

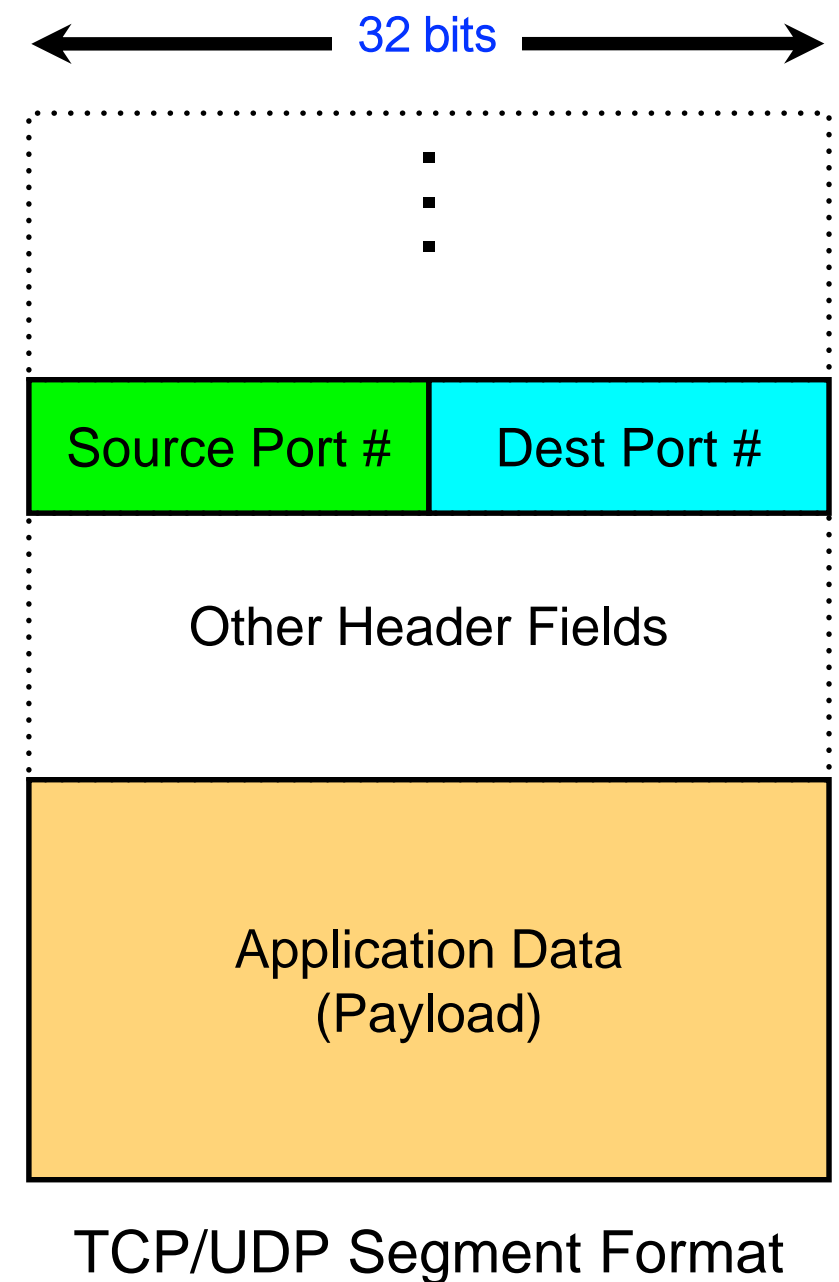
- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
 - Connectionless Demultiplexing
 - Connection-Oriented Demultiplexing
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
- **Principles of Congestion Control**
- **TCP Congestion Control**

Multiplexing/Demultiplexing



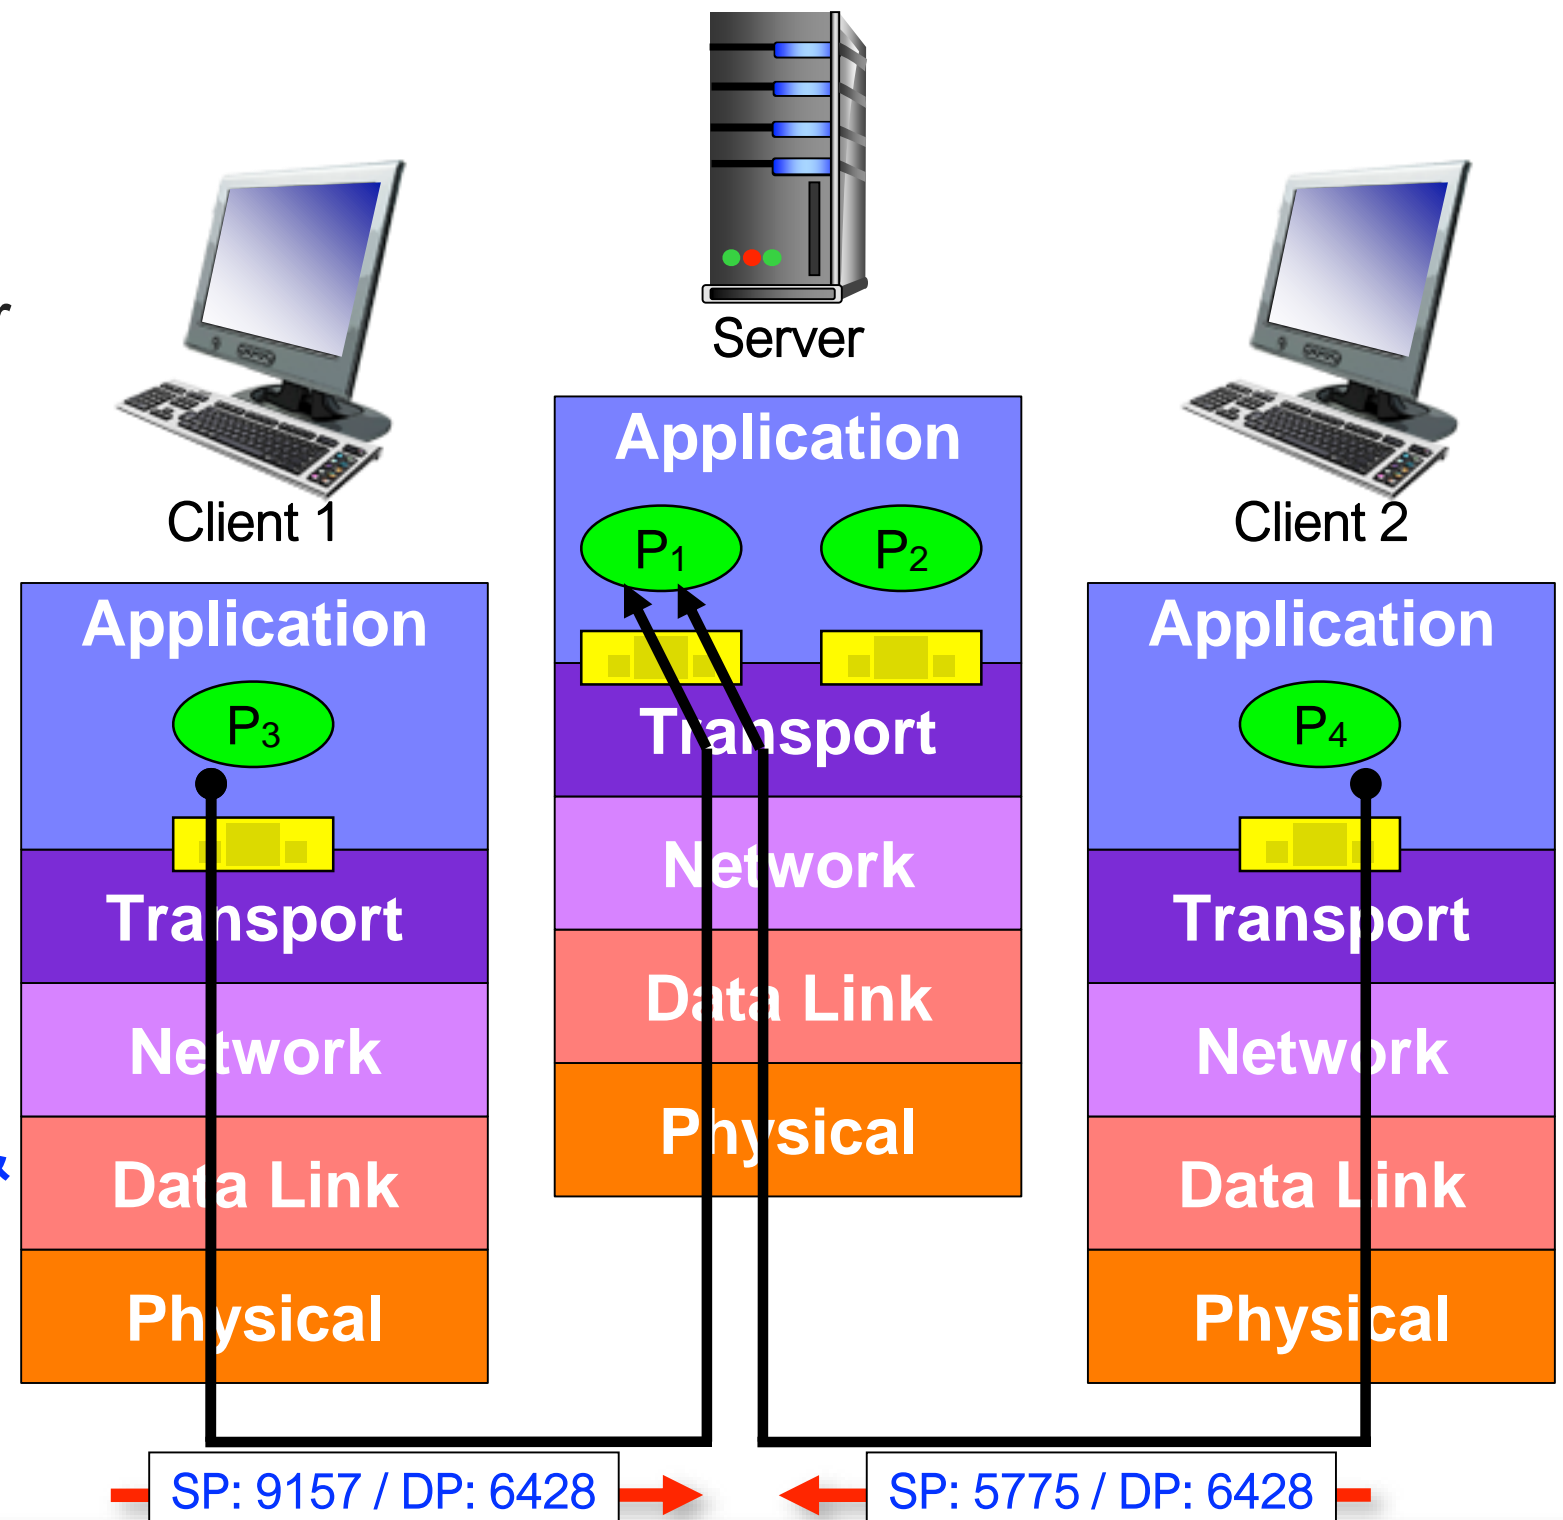
How Demultiplexing Works

- **Host receives IP datagrams**
 - Each datagram has source IP address and destination IP address
 - Each datagram carries one transport-layer segment
 - Each segment has source and destination port number
- **Host uses IP addresses & port numbers to direct segment to appropriate socket**



Connectionless Demux: Example (UDP)

- **Server UDP socket has a local port #**
 - Same socket is shared for incoming connections destined for that port #
- **Each client:**
 - Creates own local socket with own local port #
 - When sending UDP datagram to server, client must specify **IP address & port #** of server's UDP socket



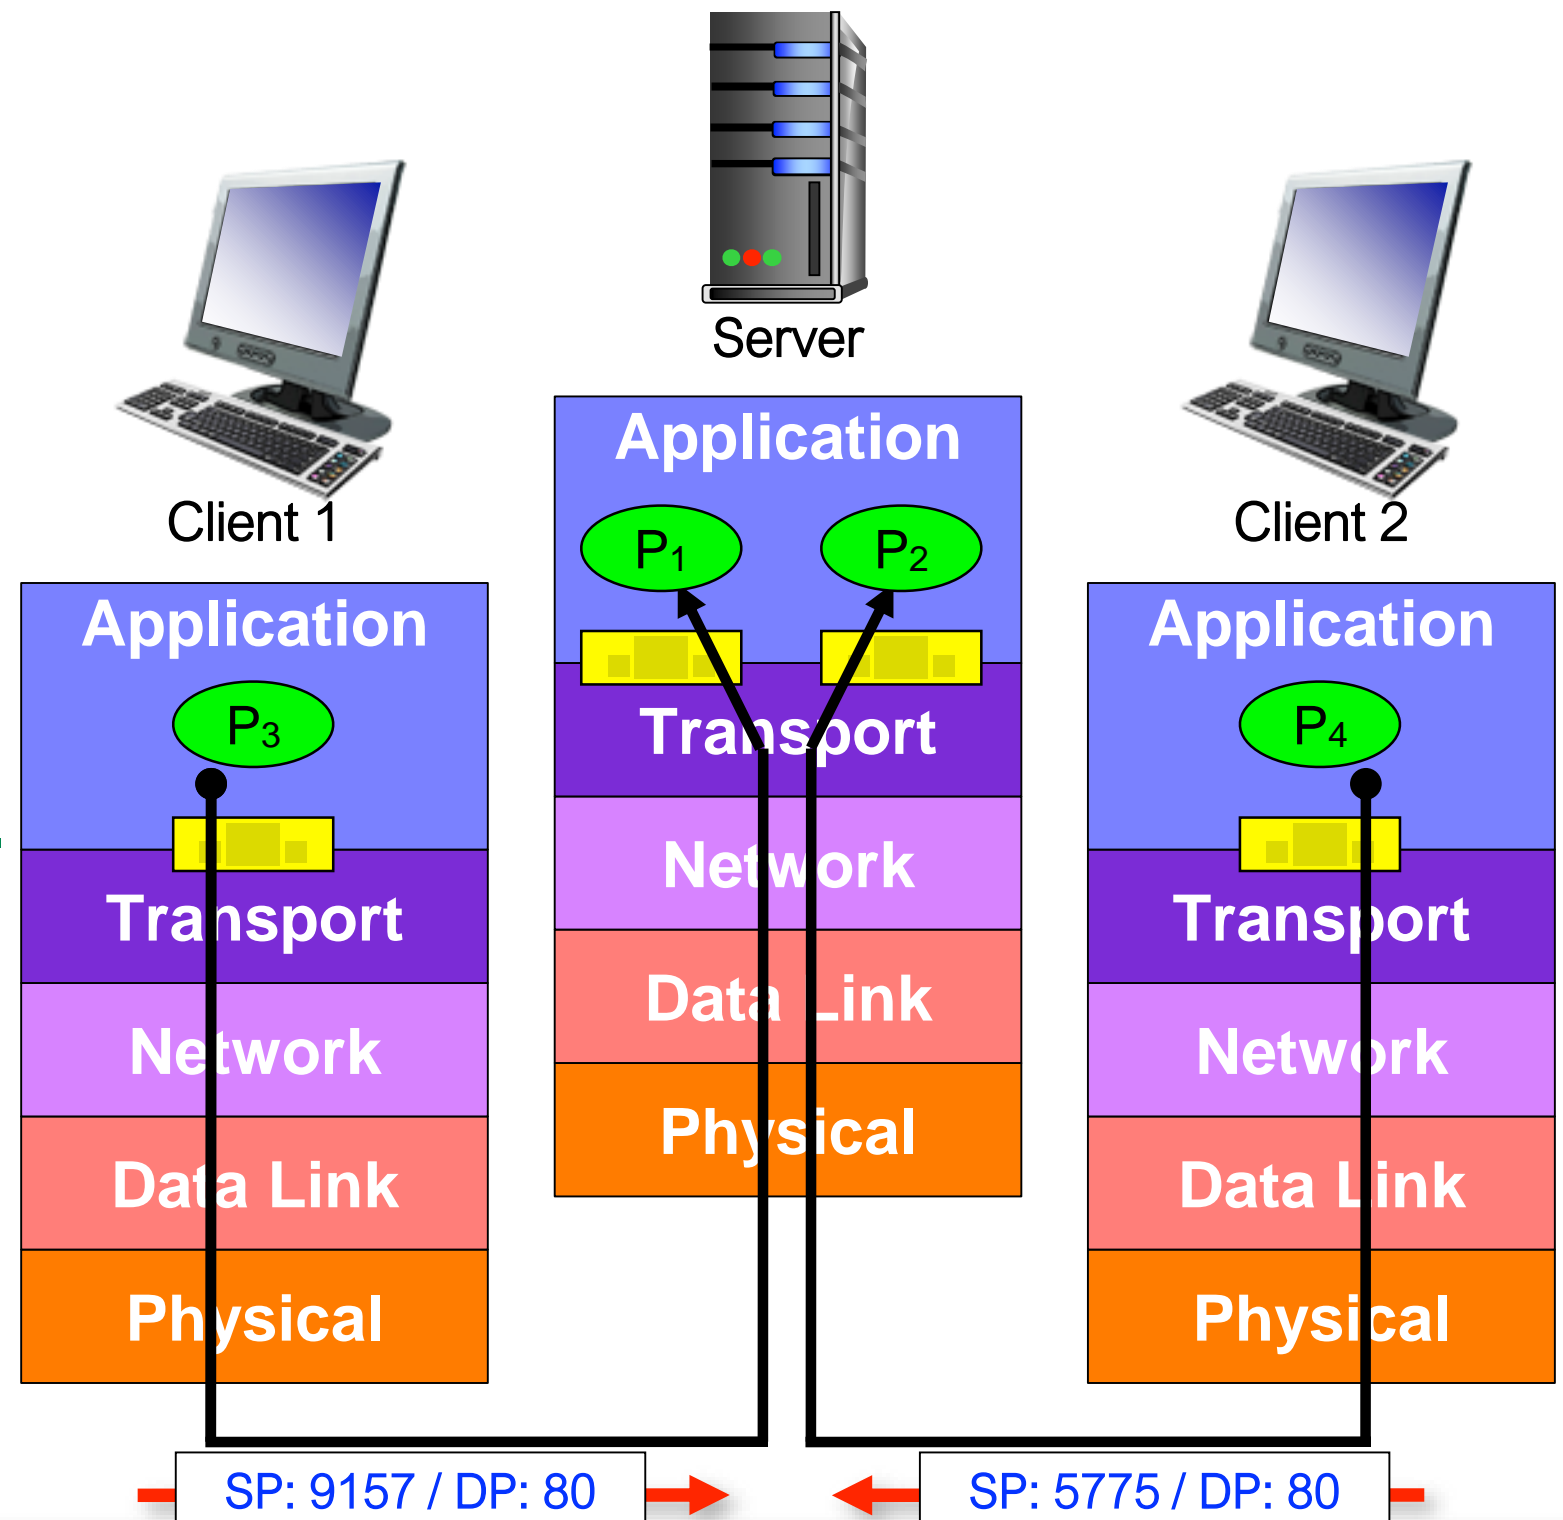
Connection-Oriented Demux: Example (TCP)

- **TCP socket is identified by a 4-tuple**

- Source IP address, Source port number, Dest IP address, Dest port number

- **A new TCP socket is created for each unique 4-tuple**

- Server host may support many simultaneous TCP sockets (even multiple from same client)




Connection-Oriented Demux: Example (TCP)

- **TCP socket is identified by a 4-tuple**

- Source IP address, Source port number, Dest IP address, Dest port number

- **A new TCP socket is created for each unique 4-tuple**

- Server host may support many simultaneous TCP sockets (even multiple from same client)



Overview of Transport Layer

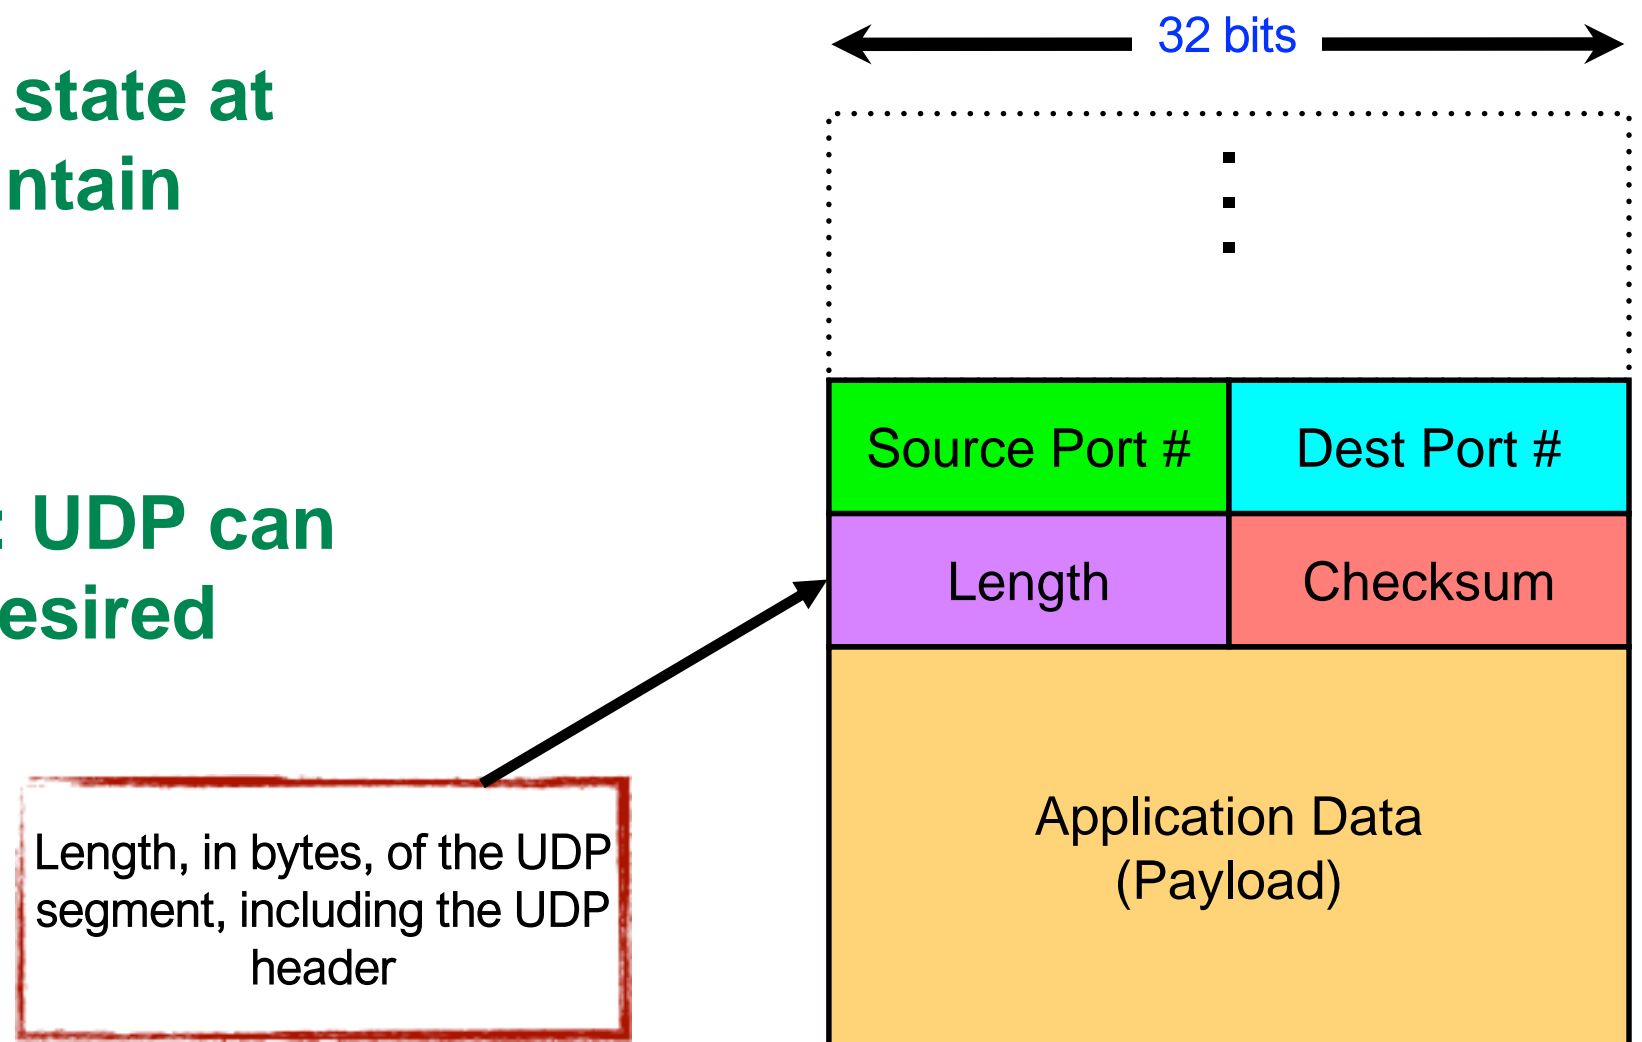
- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
 - Overview
 - Checksum
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
- **Principles of Congestion Control**
- **TCP Congestion Control**

UDP: User Datagram Protocol [RFC 768]

- **“No frills” / “bare bones” Internet transport protocol**
- **Best effort service, UDP segments may be:**
 - Lost
 - Delivered out-of-order to application
- **Connectionless:**
 - **No handshaking** between UDP sender and receiver
 - Each UDP segment is handled independently of others
- **UDP use:**
 - Streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS (Domain Name System)
 - SNMP (Simple Network Management Protocol)
- **Reliable transfer over UDP:**
 - Must add reliability at application layer
 - Application-specific error recovery!

UDP: User Datagram Protocol (Cont.)

- **No connection establishment**
 - Eliminates a source of delay
- **Simple: no connection state at sender, receiver to maintain**
- **Small header size**
- **No congestion control: UDP can blast away as fast as desired**



UDP Segment Format

UDP Checksum

- **Used to detect errors (e.g. flipped bits) in transmitted data segment**
- **Sender:**
 - Treat segment contents, including the header fields, as sequence of 16-bit integers
 - Perform one's complement sum of segment contents, then take one's complement of that sum
 - Insert checksum value into UDP checksum field
- **Receiver:**
 - Compute one's complement sum of received segment (including checksum field)
 - Check if computed sum equals **0xFFFF**
 - YES - no error detected. But may have errors nonetheless? More later
 - NO - error detected

Checksum: Example

Example: add two 16-bit integers (a UDP checksum would add many more 16-bit integers)

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
Partial Sum	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
	<hr/>															
Sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
Checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

- In one's complement addition, add overflow back into the partial sum to get the sum
- Take one's complement (invert) of sum to get the checksum

Overview of Transport Layer

- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
 - Overview
 - Pipelined Protocols
 - Go-Back-N
 - Selective Repeat
- **Connection-oriented Transport: TCP**
- **Principles of Congestion Control**
- **TCP Congestion Control**

Principles of Reliable Data Transfer

- **Reliable data transfer is very important application, transport, and link layers**
- **The characteristics of unreliable channel will determine the complexity of a **Reliable Data Transfer (RDT) protocol****
- **To explore reliable data transfer, examine different types of loss and how to address them**

RDT 1.0: Reliable Transfer Over a Reliable Channel

- **Underlying data transmission channel is perfectly reliable**
 - No bit errors
 - No loss of packets
- **We don't have networks like this, but it's a good place to start**

RDT 2.0: Data Channel With Bit Errors

- **Underlying data channel may flip bits in packet**
 - Use a checksum to detect bit errors
- **Question: How should system recover from these errors?**
 - **Acknowledgements (ACKs)**: receiver explicitly tells sender that a packet is received OK
 - **Negative acknowledgements (NAKs)**: receiver explicitly tells sender that a packet had errors when it was received
 - **Sender retransmits packet on receipt of a NAK**
- **New mechanisms in RDT 2.0 (beyond RDT 1.0):**
 - Error detection
 - Feedback: send control messages (ACK, NAK) from receiver to sender

A Fatal Flaw in RDT 2.0, on to RDT 2.1

- **What happens if ACK/NAK messages are corrupted?**
 - Sender doesn't know if receiver correctly received the data
 - Data may have been corrupt on the way to receiver
 - ACK/NAK may have been corrupt on the way back to sender
 - Can't just retransmit since receiver may receive duplicate data!
- **Handling duplicates:**
 - Sender retransmits current packet if ACK/NAK is corrupted
 - Sender adds a **sequence number** to each packet
 - Receiver discards duplicate packets at Transport Layer
 - Those packets are not delivered up to the Application Layer

RDT 2.1: Discussion

- **Sender:**

- Sequence number added to packets
- Two sequence numbers (0,1) will suffice
- Must check if received ACK/NAK corrupted
- Sender must “remember” if it should be expecting a sequence number of 0 or 1

- **Receiver:**

- Must check if received packet is duplicate
 - Receiver must “remember” if it should be expecting a sequence number of 0 or 1
- Note: the receiver does not know if its last ACK/NAK message was received OK at sender


RDT 2.2: Eliminating the NAK Messages

- **Possible to achieve same functionality as RDT 2.1 using ACKs messages only**
- **Receiver sends ACK messages for last packet that was received correctly**
 - No message is sent for a packet this is received with errors
- **Duplicate ACK messages received at sender results in same action as NAK from RDT 2.1 -- retransmit packet**
 - Duplicate ACK message would be detected at sender by receiving two consecutive ACK_0 messages or two consecutive ACK_0 messages


RDT 3.0: Channels with Bit Errors & Packet Loss

- **New assumption: underlying channel can also lose packets**
 - Can lose data packets or ACK messages
 - Checksum, sequence number, ACKs, retransmissions will be of help ... but not enough
- **Approach: sender waits a “reasonable” amount of time for an ACK**
 - Retransmits packet if no ACK is received in this time
 - If packet (or ACK) is just delayed and not lost:
 - Retransmission will result in duplicate data, but sequence numbers from RDT 2.2. already handle this issue
 - Receiver must specify sequence number of packet being ACKed
 - Requires countdown timer

RDT 3.0 in Action with Packet Loss

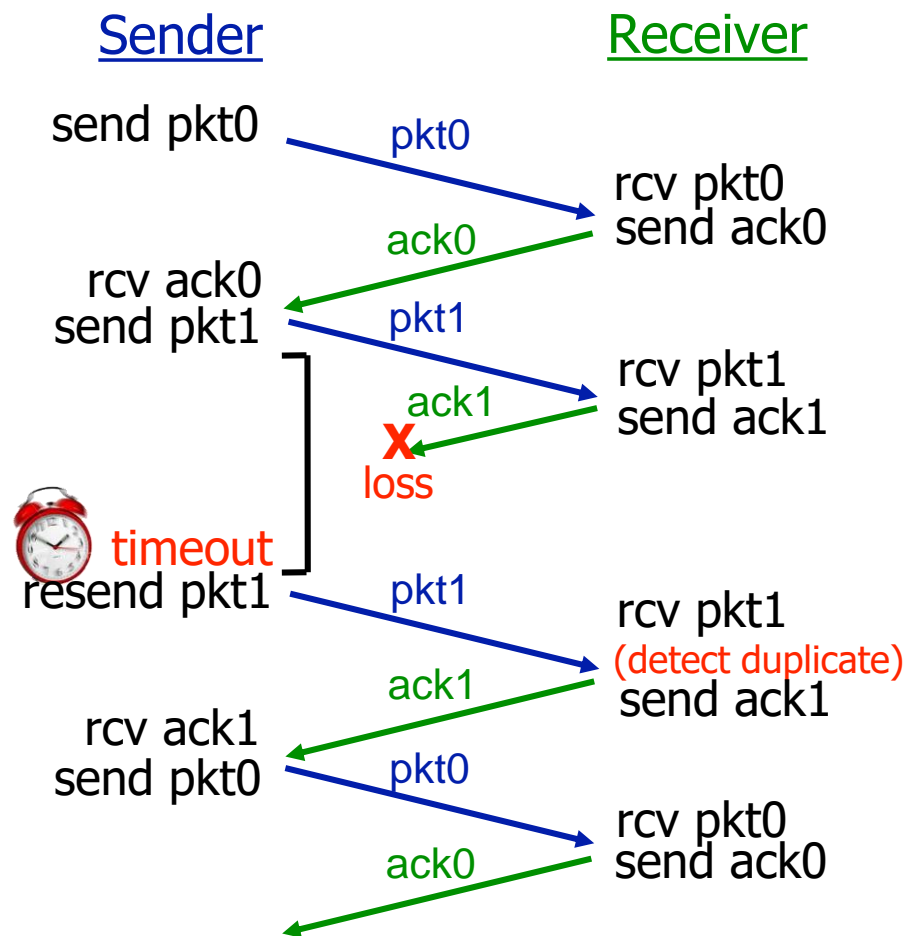


No Packet Loss




With Packet Loss

RDT 3.0 in Action with Lost/Delayed ACK



ACK Loss



Premature Timeout/Delayed ACK

Performance of RDT 3.0

- **RDT 3.0 will work reliably, but would have terrible performance**
 - RDT 3.0 utilizes a **Stop-and-Wait protocol**
 - Sender sends one packet, then waits for receiver response before sending the next packet
- **Example:** Assume a system sending 8000 bit packets on a 1 Gbps link, where there is a 15 ms propagation delay between the sender and the receiver

$$\text{Time for sender to transmit data} \quad D_{\text{trans}} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microsecs}$$

Performance of RDT 3.0 (Cont.)

- **If Round-Trip Time (RTT) is 30 ms:**


- Sender is only sending 8 microseconds
- Can only send a new packet every 30.008 microseconds
- Effectively makes 1 Gbps link run at ~270 Kbps!!!

Utilization: Fraction
of time the sender
is busy sending

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$


.027%

The network protocol
limits use of physical
resources!

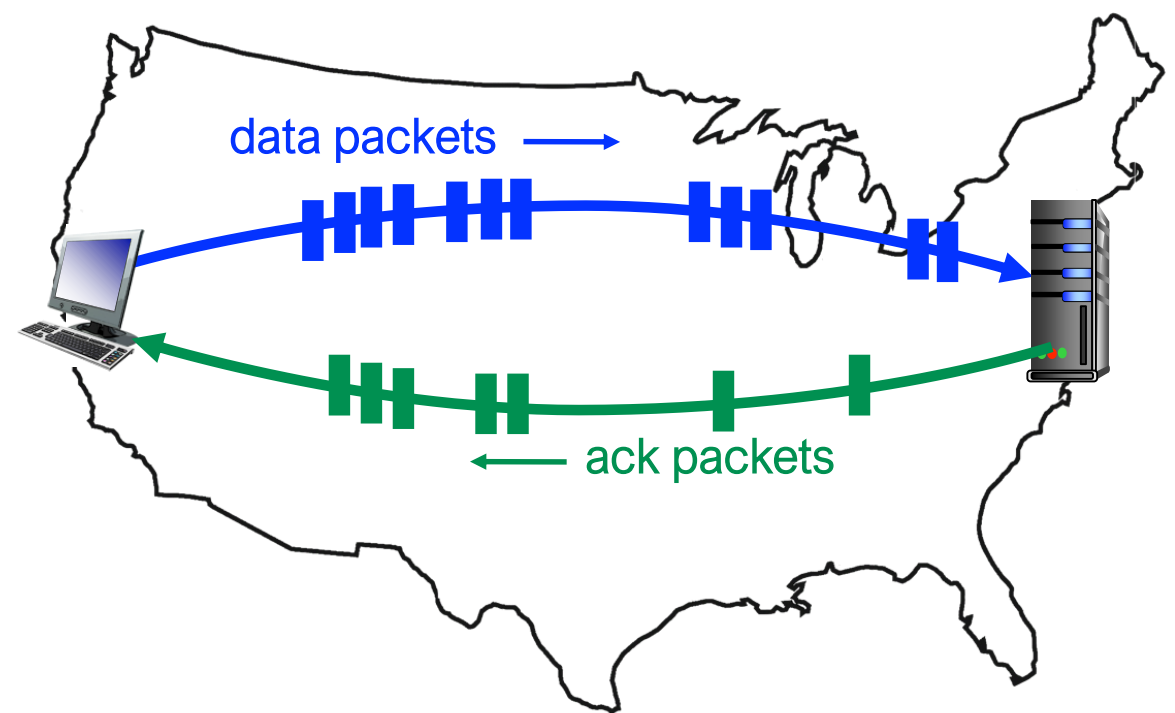


Pipelined Protocols

- **Pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged packets
 - Range of sequence numbers must be increased (0 and 1 will no longer suffice)
 - Buffering at sender and/or receiver is required
- **Two generic forms of pipelined protocols: Go-Back-N, and Selective Repeat**




Stop-and-Wait Protocol



Go-Back-N Protocol

Pipelining: Increased Utilization



Utilization: Fraction of time the sender is busy sending

$$U_{\text{sender}} = \frac{3L/R}{RTT + L/R} = \frac{.024}{30.008} = 0.0008$$

Pipelined Protocols: Overview

- **Go-back-N:**


- Sender can have up to N unacked packets in pipeline
- Receiver only sends cumulative acks
 - Doesn't ack a packet if there is a gap
- Sender has a timer for the oldest unacked packet
 - When timer expires, retransmit **all** unacked packets

- **Selective Repeat:**

- Sender can have up to N unacked packets in pipeline
- Receiver sends individual acks for each packet
- Sender maintains timer for each unacked packet
 - When timer expires, retransmit only that unacked packet

Go-Back-N: Sender Side

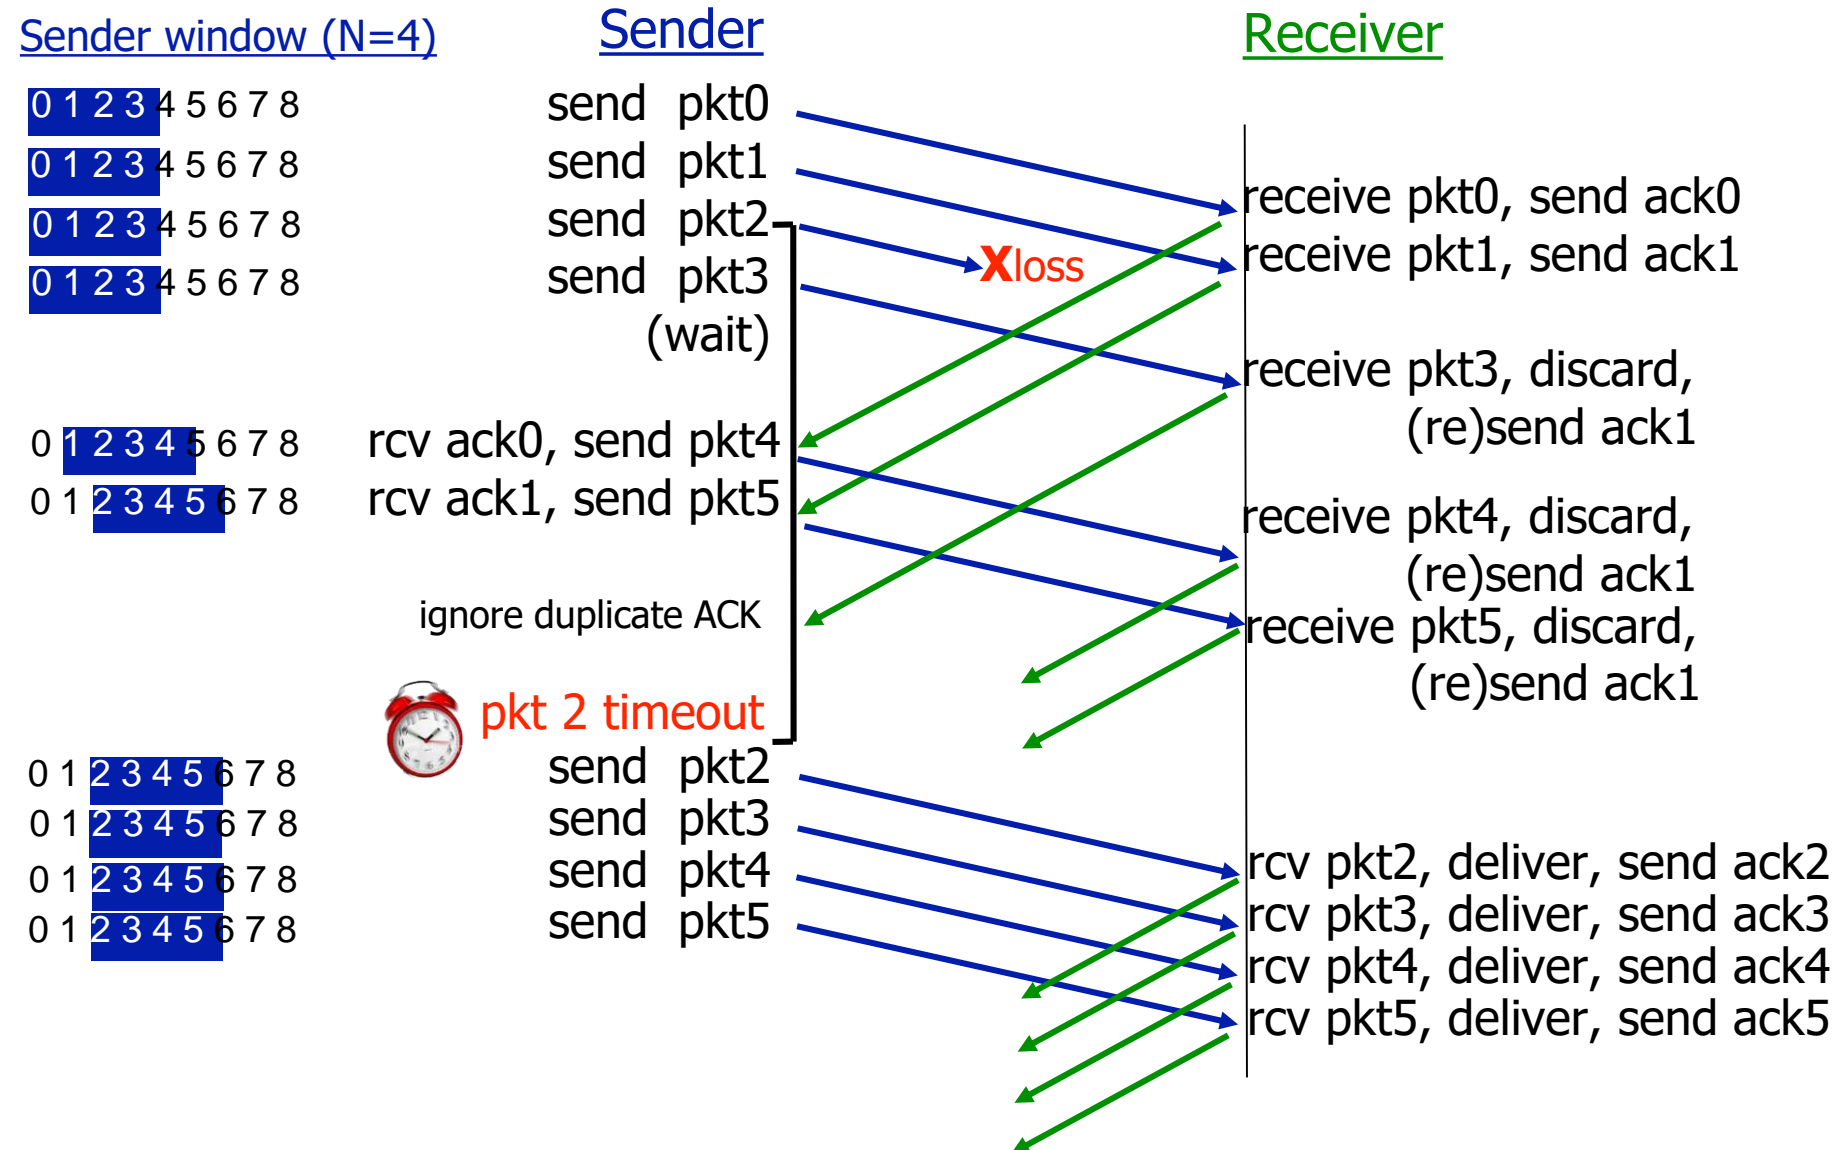
- **Sender can transmit multiple packets without waiting for an ack**
 - A “sliding window” of up to N consecutive unacked packets allowed
- **Include a k-bit sequence number in the packet header**
- **ACK(n): ACKs all packets up to and including sequence number n**
 - Cumulative ACK
 - May receive duplicate ACKs
- **Maintain timer for oldest in-flight packet**
- **Timeout(n): retransmit packet n and all higher sequence number packets in window**



Go-Back-N: Receiver

- **Send ACK message for correctly-received packet with highest in-order sequence number**
 - May generate duplicate ACKs
 - Need only remember expected sequence number
 - Must receive packets in-order to send ACK
- **Out-of-order packets:**
 - **Discard** (don't buffer) out-of-order packets (no receiver buffering)
 - Yes, even if they are correctly formatted and error-free
 - Will be retransmitted anyway based on sender's rules
 - Re-ACK packet with highest in-order sequence number

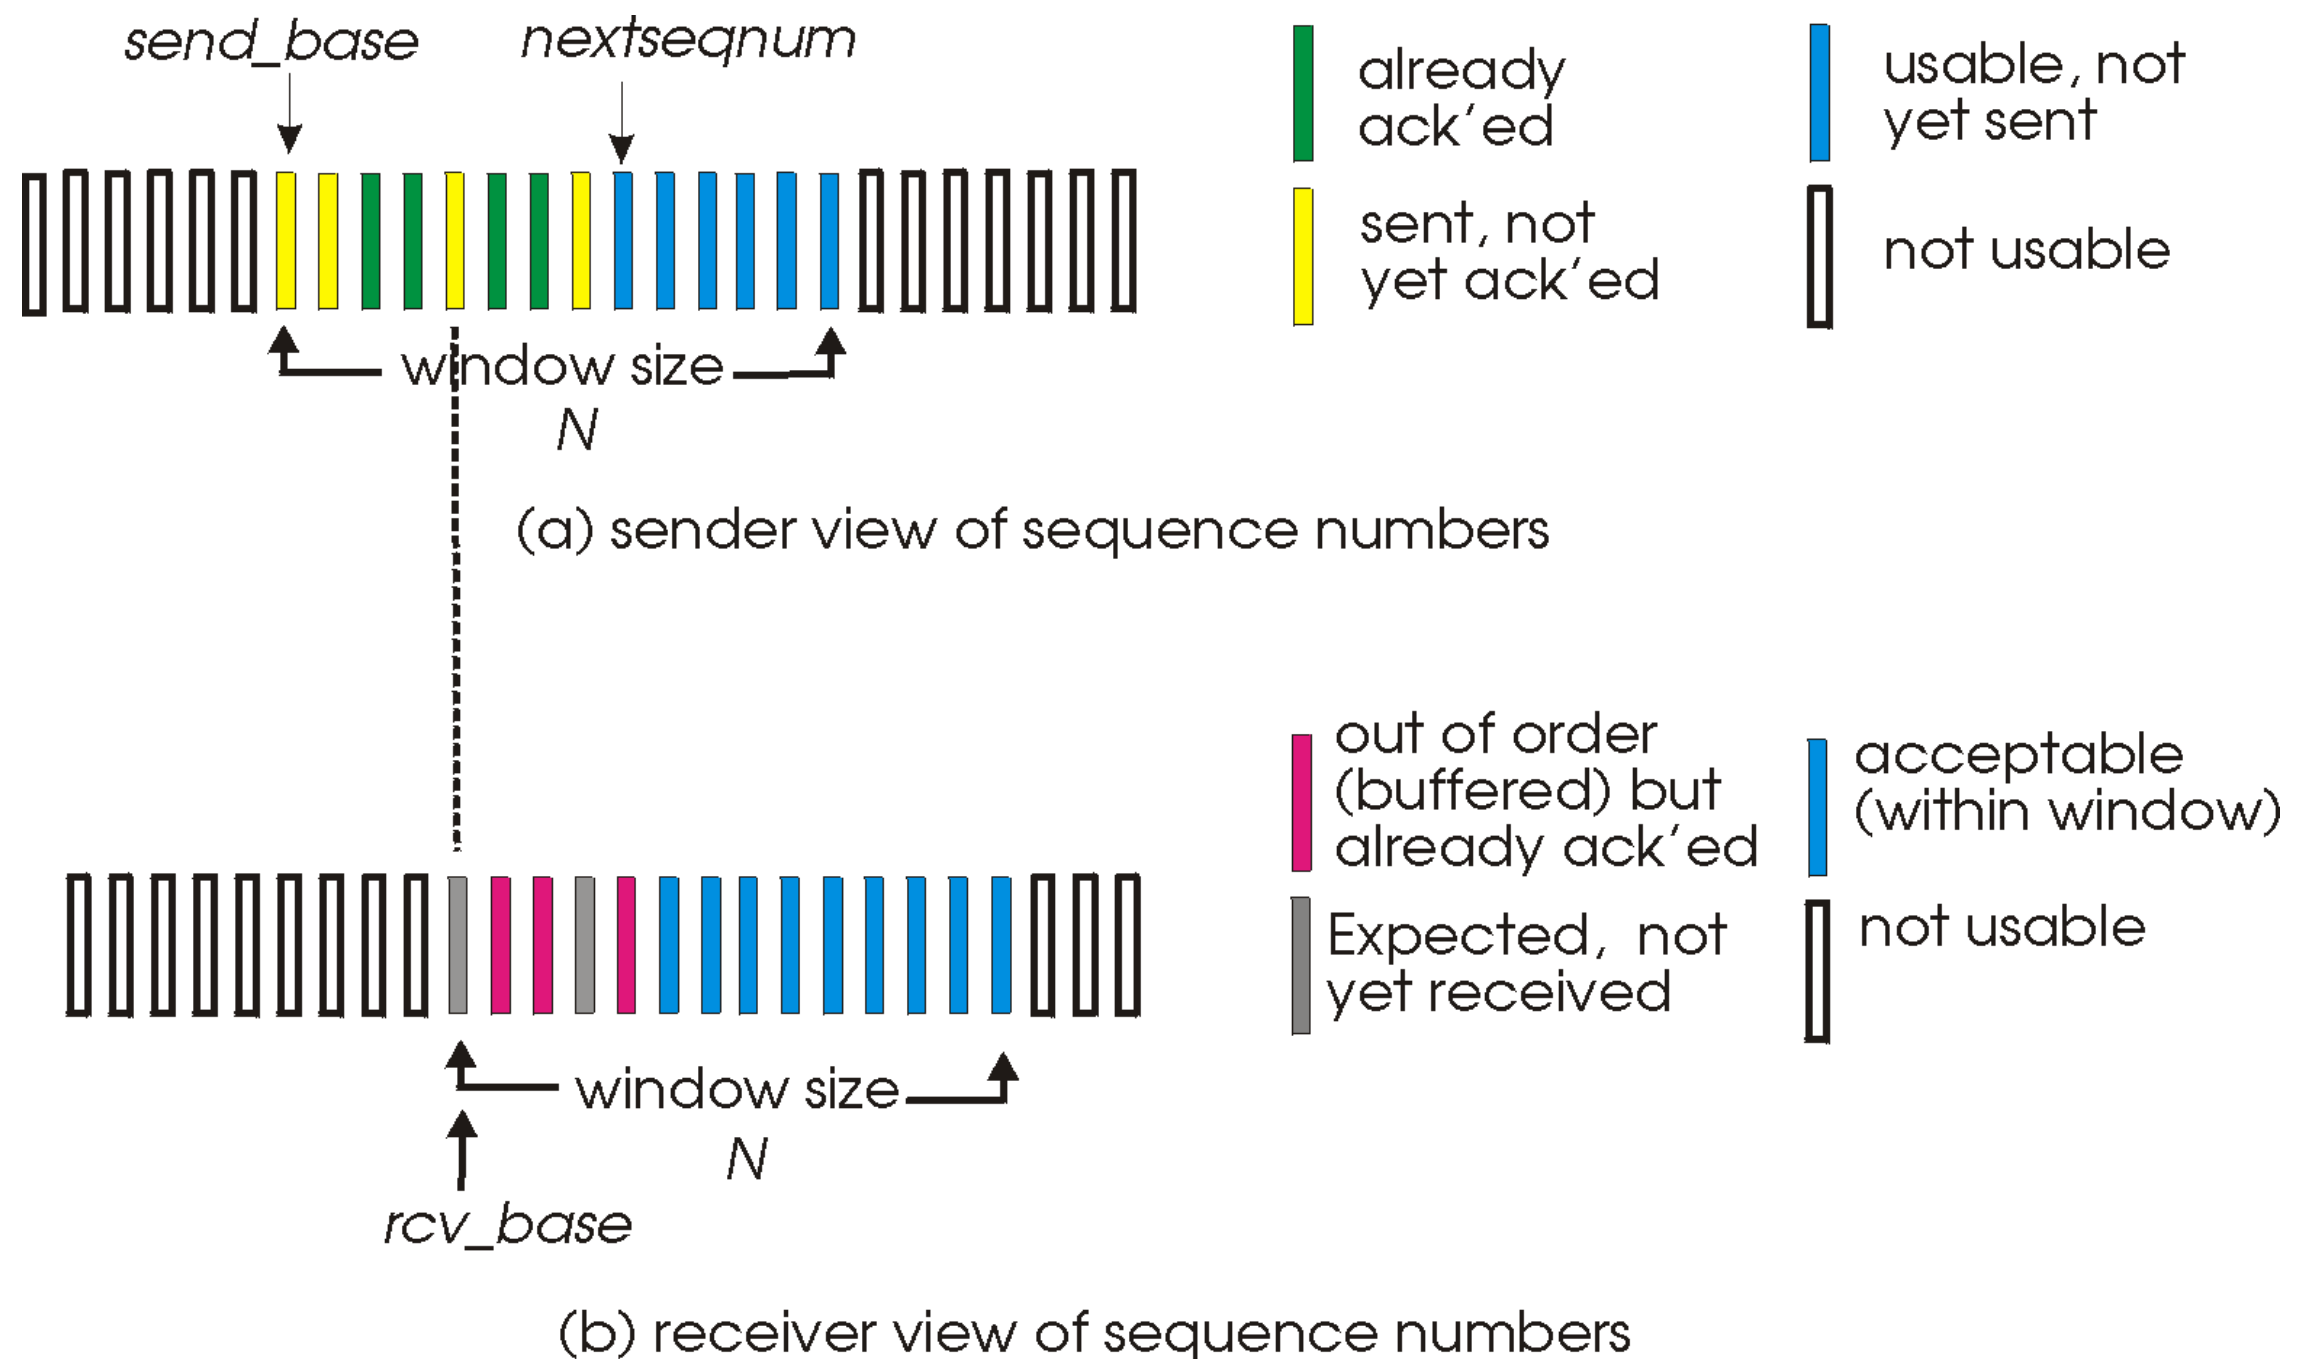
Go-Back-N in Action



Selective Repeat

- **Receiver individually acknowledges all correctly received packets**
 - Out-of-order packets are buffered at the receiver
 - Buffered packets are eventually delivered in-order to upper layer
- **Avoids unnecessary retransmission -- sender only resends packets for which an ACK was not received**
 - Sender has separate timer for each unACKed packet
- **Sender window**
 - N consecutive sequence numbers
 - Limits sequence numbers of sent, unACKed packets

Selective Repeat: Sender/Receiver Windows



Selective Repeat

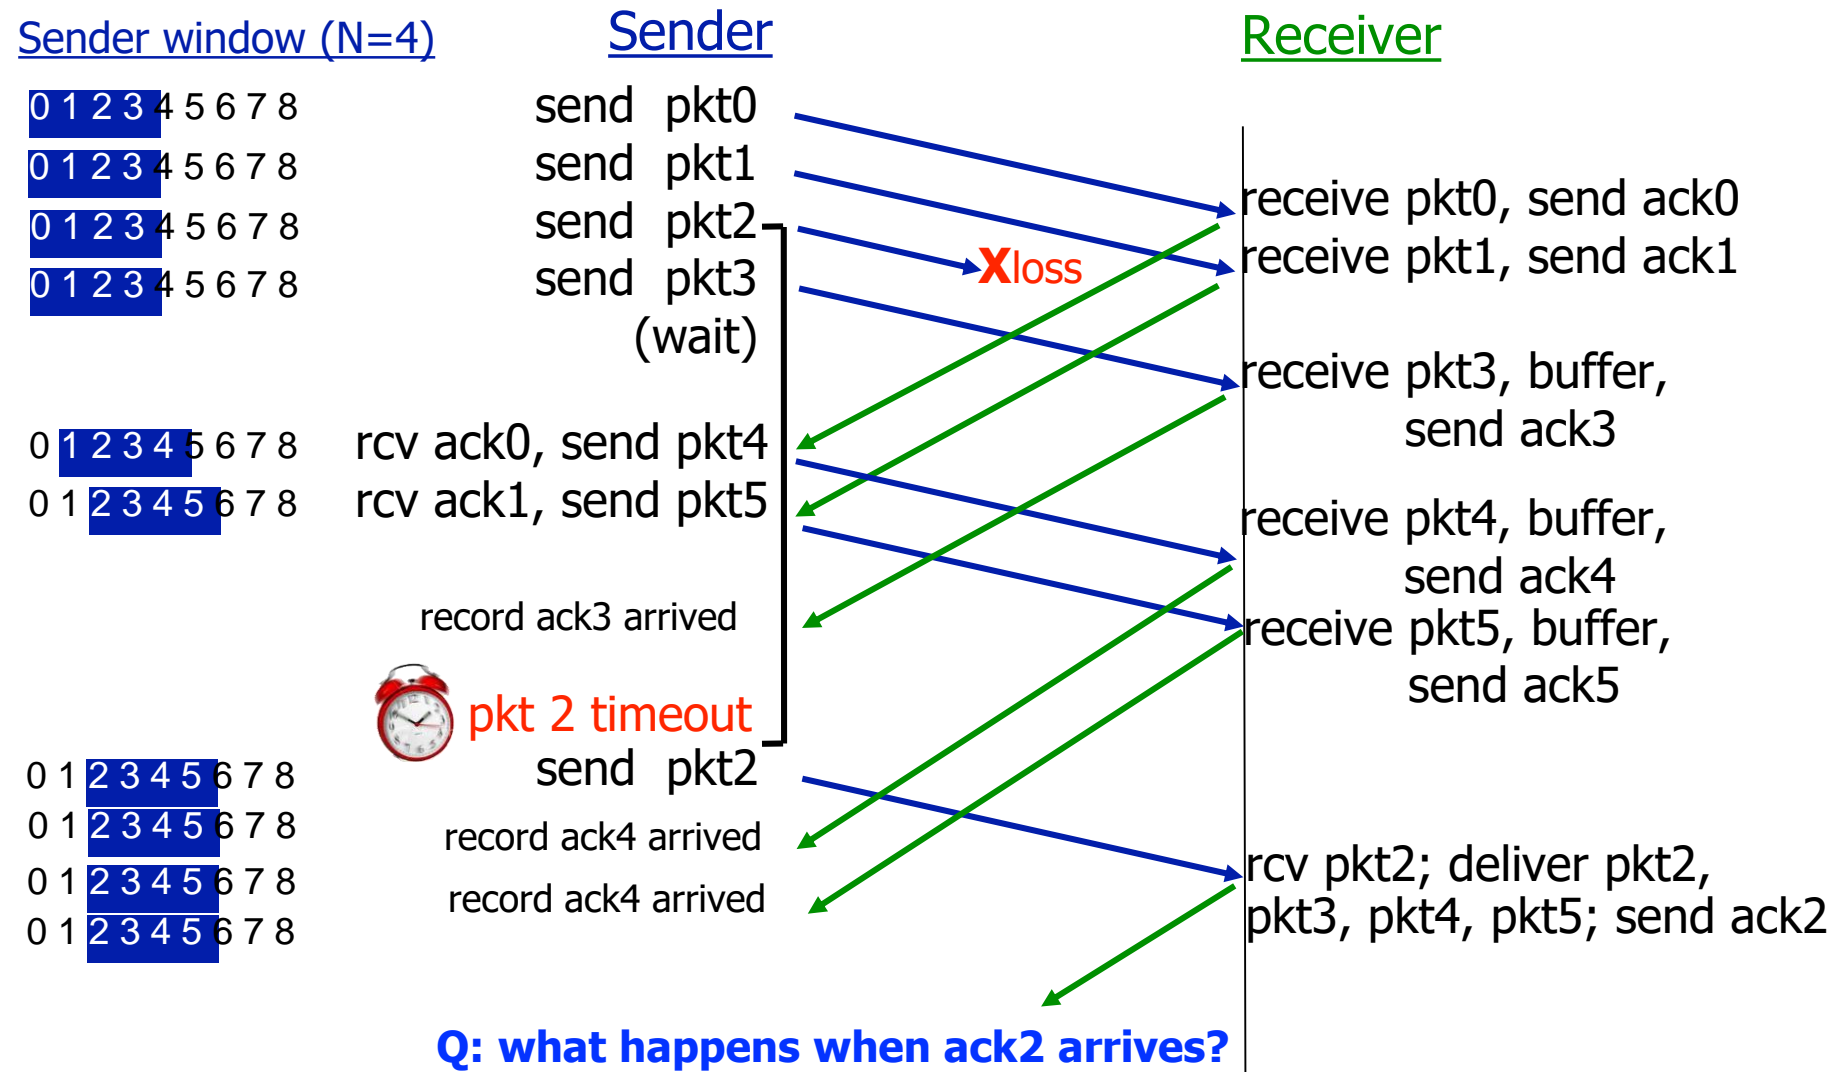
- **Sender**

- Receives data from upper layer:
 - If next available sequence number is in the sliding window, send the packet
- A timeout occurs for packet n:
 - Resend packet n, restart timer
- ACK(n) received for packet in current window:
 - Mark packet n as received
 - If n is smallest unACKed packet in window, advance the window base to next unACKed sequence number

- **Receiver**

- Receives packet n in receivers window
 - Send ACK(n)
 - If out-of-order, buffer
 - If in-order, deliver with other buffered, in-order packets to the upper layer. Also, advance window to next not-yet-received packet
- Receives packet n that has already been seen and ACKed by the receiver
 - Send ACK(n) again
- Otherwise:
 - Ignore the packet

Selective Repeat in Action



Overview of Transport Layer

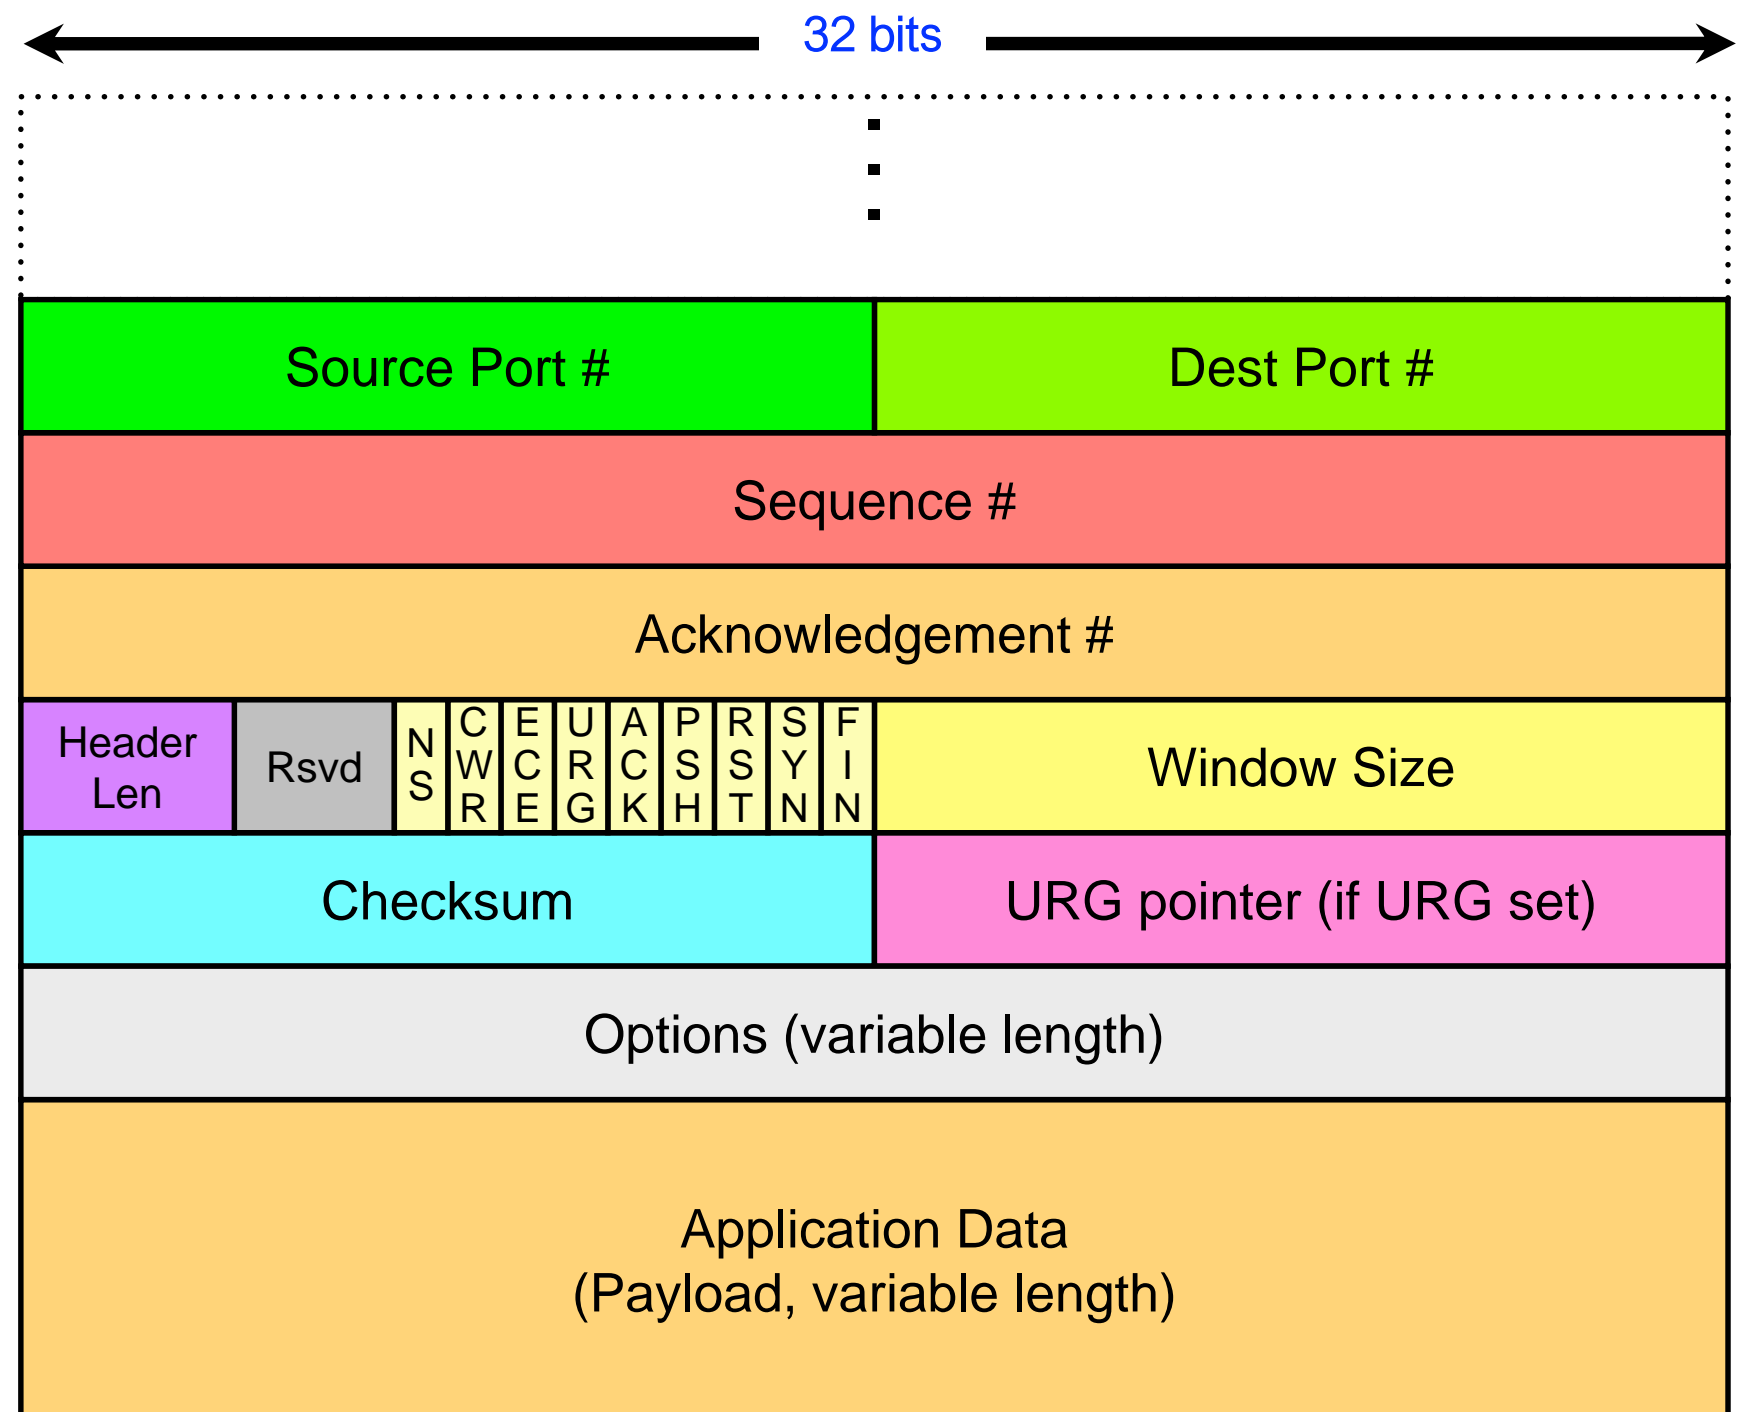
- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
 - Segment Structure
 - Reliable Data Transfer
 - Flow Control
 - Connection Management
- **Principles of Congestion Control**
- **TCP Congestion Control**

TCP: Overview

- **A point-to-point protocol - one sender, one receiver**
- **Provides a reliable, in-order byte stream**
- **A pipelined protocol -- allows multiple in-flight packets**
 - TCP congestion and flow control set window size
- **Full duplex data - bi-directional data flow on same connection**
- **Connection-oriented:**
 - Handshaking (the exchange of control messages) initializes sender and receiver state before data exchange starts
- **Flow controlled:**
 - Sender will not overwhelm receiver


TCP Segment Structure

- **Source Port #** - the port on the sender
- **Dest Port #** - the port on the receiver
- **Sequence #** - 32-bit number that represents the byte stream 'number' of the first byte in this segment's data
 - e.g. if Sequence # is 10 and there are 7 bytes of data in this packet, then this segment contains bytes 10-16 of the data stream
- **Acknowledgement #** - 32-bit number that represents the next sequence number that the receiver is expecting



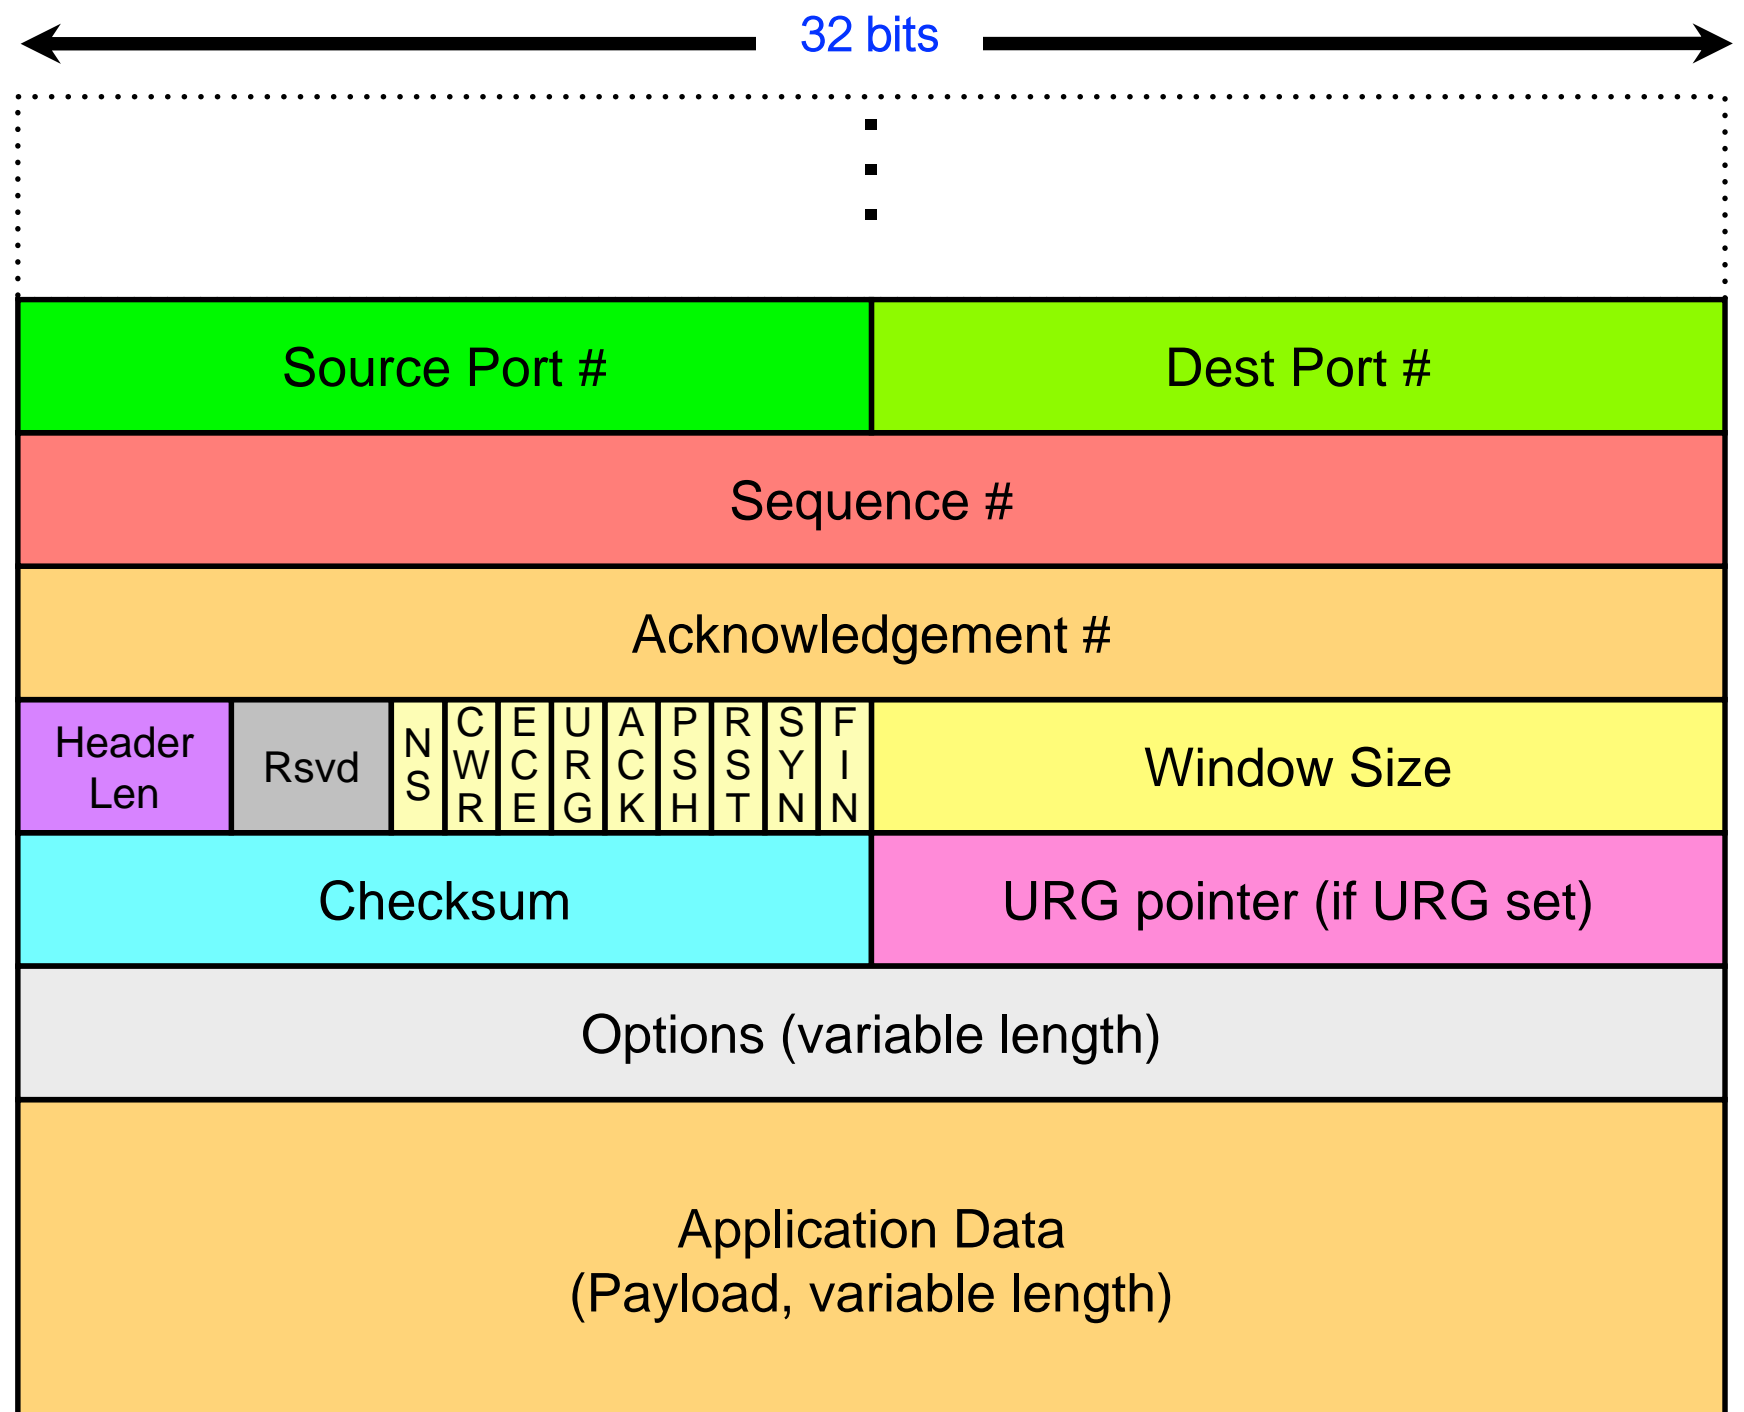
TCP Segment Structure (Cont.)

- **Header Length** - 4-bit field that specifies the size of the TCP header in 32-bit words
 - min = 5 ; max = 16
- **Reserved** - 3-bits, not currently used
- **Window Size** - the size of the receive window that specifies the number of bytes that the receiver of this packet is currently willing to receive
- **URG Pointer** - specifies an offset of urgent data
- **Checksum** - 16-bit checksum computed over header and data (similar to UDP)
- **Options** - optional header fields



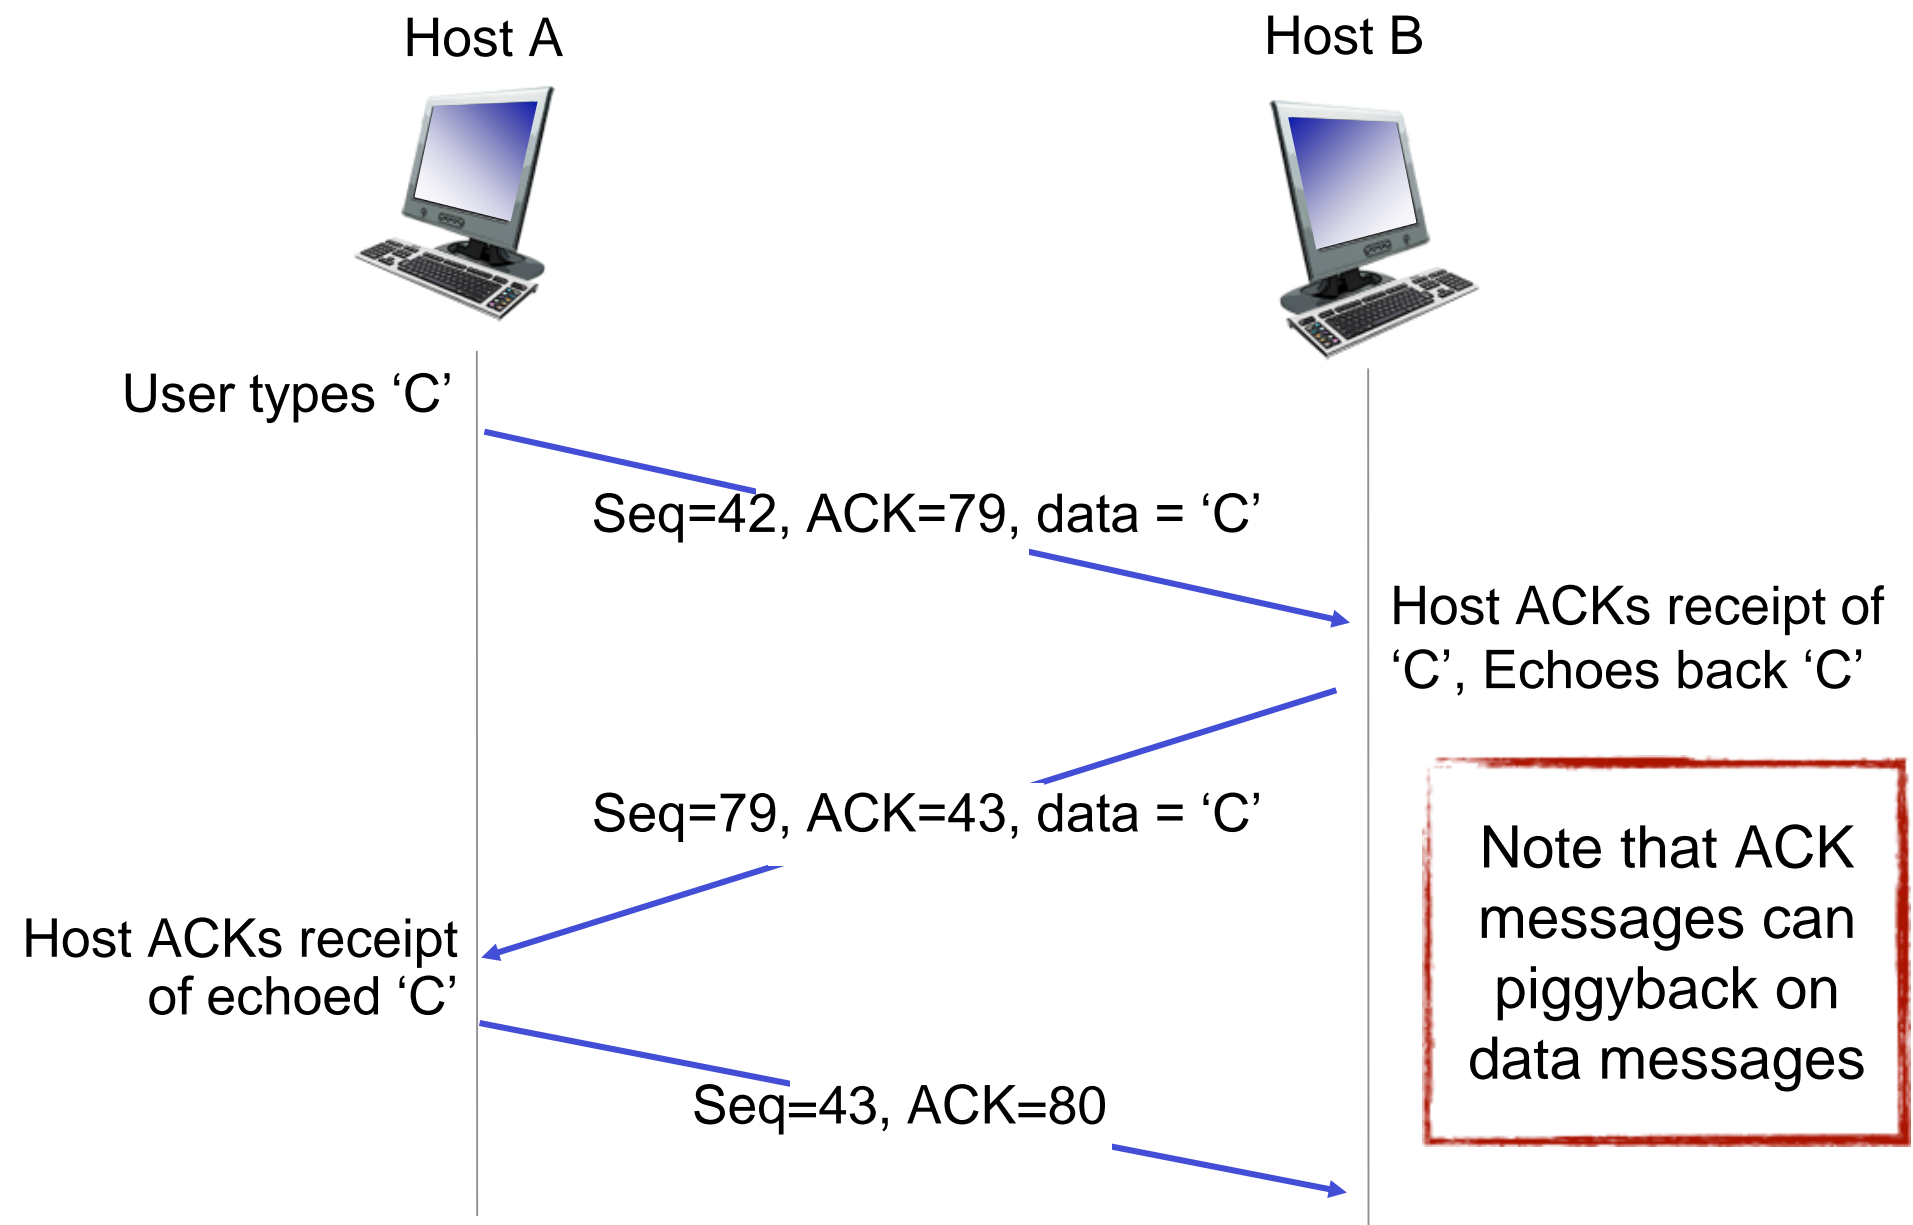
TCP Segment Structure: Flags (1-bit each)

- **NS, CWR, & ECE** - used in congestion mechanism
- **URG** - indicates that an URG pointer is present (not used much)
- **ACK** - indicates that the ACK # is significant
- **PSH** - request to push data to receiving application (not used much)
- **RST** - reset the connection
- **SYN** - synchronize sequence numbers. First packet from each end of connection should set this.
- **FIN** - indicates that there is no more data from the sender



TCP Sequence Numbers/ACKs

Simple Telnet Scenario




TCP Round Trip Time & Timeout

- **How should the TCP timeout value be set?**
 - Must be longer than RTT -- but RTT varies
 - Too short and timer will timeout prematurely causing unnecessary retransmissions of data
 - Too long and sender will have a slow reaction to segment loss
- **How can the RTT be estimated?**
 - Sample the RTT -- measure the time from segment transmission until ACK receipt (**SampleRTT**)
 - Ignore retransmissions
 - **SampleRTT** will vary, need to compute average over time to 'smooth' value

TCP Round Trip Time & Timeout

- Compute EstimatedRTT from SampleRTT
- Exponentially weighted moving average
- Influence of past sample decreases exponentially fast
- Typical value: $\alpha = 0.125$
- Use EstimatedRTT plus some safety-margin to determine timeout interval

$$\text{EstimatedRTT} = (1-\alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$



Overview of Transport Layer

- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
 - Segment Structure
 - **Reliable Data Transfer**
 - Flow Control
 - Connection Management
- **Principles of Congestion Control**
- **TCP Congestion Control**

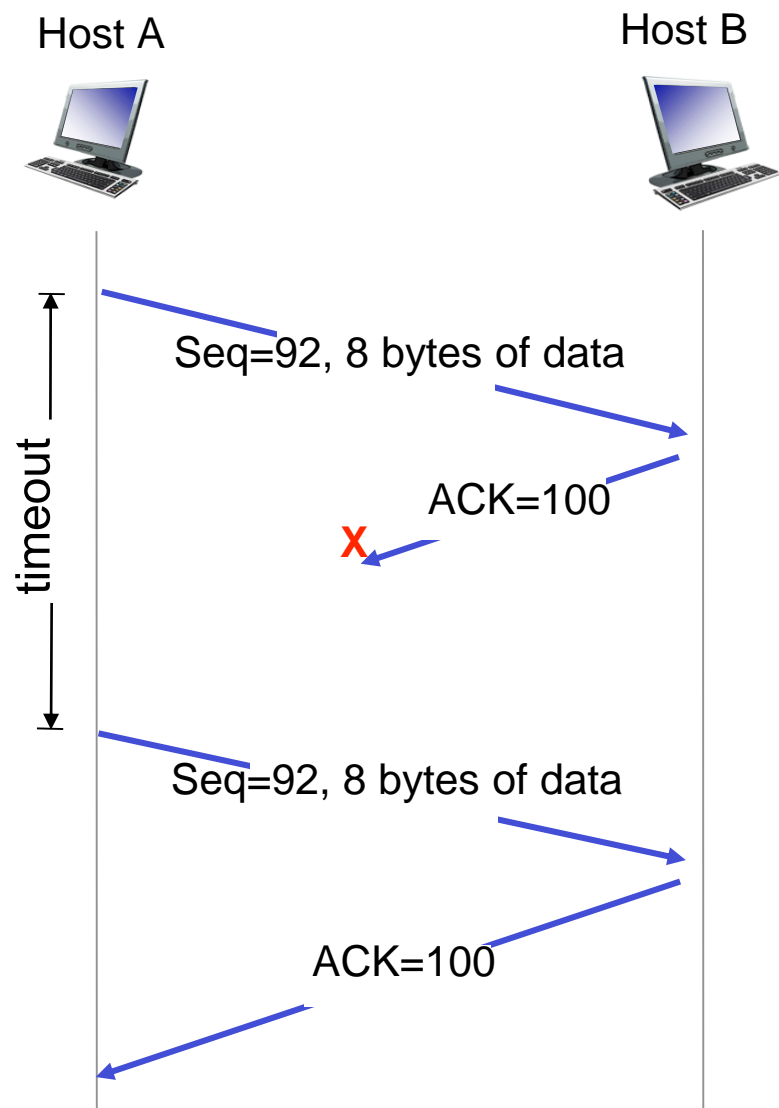
TCP Reliable Data Transfer

- **TCP creates reliable data transfer service on top of the IP protocol's unreliable service**
 - Send data as pipelined segments
 - Uses cumulative ACKs
 - Uses a single retransmission timer
- **Retransmissions are triggered by:**
 - Timeout events
 - Duplicate ACK messages

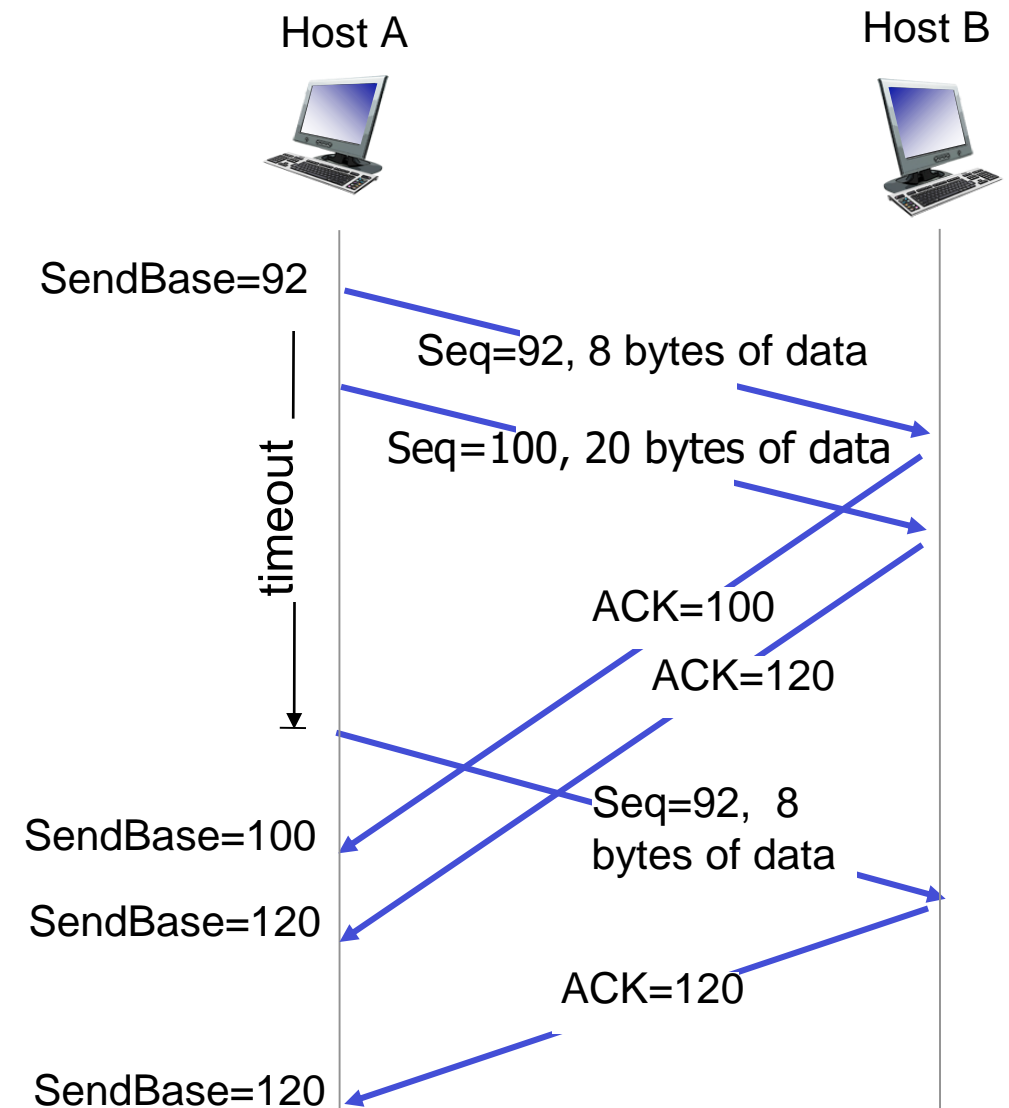
TCP Sender Events

- **Data received from application layer**
 - Create segment with sequence number
 - Sequence number is byte-stream number of the first data byte in segment
 - Start a timer if not already running
 - Timer is for oldest unacked segment
 - Use timeout interval computed from sampling RTT
- **If timeout occurs:**
 - Retransmit the segment that caused timeout
 - Restart the timer
- **When ACK is received:**
 - If ACK is for previously unacked segments
 - Update what is known to be ACKed
 - Restart timer if there are still unacked segments

TCP: Retransmission Scenarios

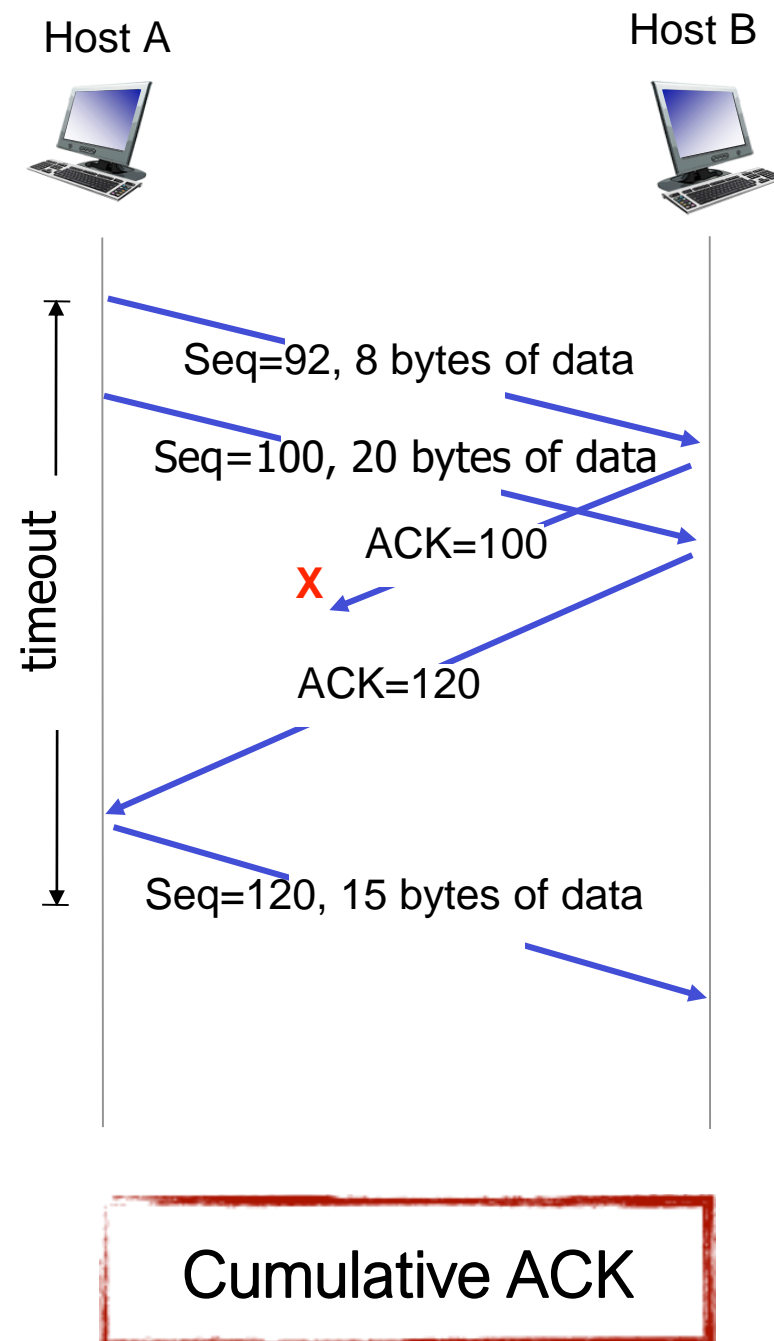


Lost ACK



Premature Timeout

TCP: Retransmission Scenarios (Cont.)




TCP ACK Generation

<i>Event at Receiver</i>	<i>TCP Receiver Action</i>
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK. Wait up to 500 ms for next segment. If no segment arrives, send ACK
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediately send a single cumulative ACK, ACKing both in-order segments
Arrival of out-of-order segment higher-than-expect seq. #. Gap detected	Immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
Arrival of a segment that partially or completely fills gap	Immediate send ACK, provided that segment starts at lower end of gap

TCP Fast Retransmit

- **Time-out period for a segment is often relatively long**
 - Means there is a long delay before sender resends lost packet
- **Lost segments can usually be detected via duplicate ACKs before timeout occurs on that segment**
 - Sender often sends many segments back-to-back
 - If a segment is lost, there will likely be many duplicate ACKs indicating that the lost segment should be retransmitted
- **TCP fast retransmit protocol**
 - If sender receives 3 ACKs for same data (**triple duplicate ACKs**), resend unacked segment with the smallest sequence number
 - **Very likely that the unacked is segment lost, so don't wait for timeout**

TCP Fast Retransmit (Cont.)



Overview of Transport Layer


- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
 - Segment Structure
 - Reliable Data Transfer
 - **Flow Control**
 - Connection Management
- **Principles of Congestion Control**
- **TCP Congestion Control**

TCP Flow Control

- **The receiver can control the sender, so the sender won't overflow the receiver's buffer by transmitting too much, too fast**
 - TCP socket between Transport and Application layer contains buffers to accumulate received data
 - Data buffer may fill faster than receiver's Application Layer app can empty the buffer
- **Receiver needs way to indicate to sender that its data buffers are too full and the sender should send data at a slower rate**

TCP Flow Control (Cont.)

- Receiver “advertises” how much buffer space it has available by including a value **rwnd**, in the **Window Size field** of the TCP header
 - The **rwnd** value indicates how much free space is available in the buffer
 - Receive buffer size can typically be set via socket options (default is usually 4096 bytes)
 - Many operating systems will automatically adjust the size of the receive buffer
- The sender limits amount of unacked (“in-flight”) data to receiver based on the **Window Size field**
 - Guarantees receive buffer will not overflow



Overview of Transport Layer

- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
 - Segment Structure
 - Reliable Data Transfer
 - Flow Control
 - Connection Management
- **Principles of Congestion Control**
- **TCP Congestion Control**

Connection Management

- **Before exchanging data, the sender and receiver perform a 3-way-handshake:**
 - Agree to establish connection (each knowing the other is willing to establish connection)
 - Agree on connection parameters

TCP 3-way Handshake


Client State

LISTEN
↓
SYN SENT
↓
ESTAB

Choose initial seq num, x
Send TCP SYN message



SYN bit=1, Seq=x



Choose initial seq num, y
Send TCP SYNACK
message to ACK the SYN

SYN bit=1, Seq=y
ACK bit=1; ACK num=x+1

Received SYNACK(x)
indicating server is live;
Send ACK for SYNACK;
This segment may contain
client-to-server data

ACK bit=1, ACK num=y+1

Received ACK(y)
indicating client is live

Server State

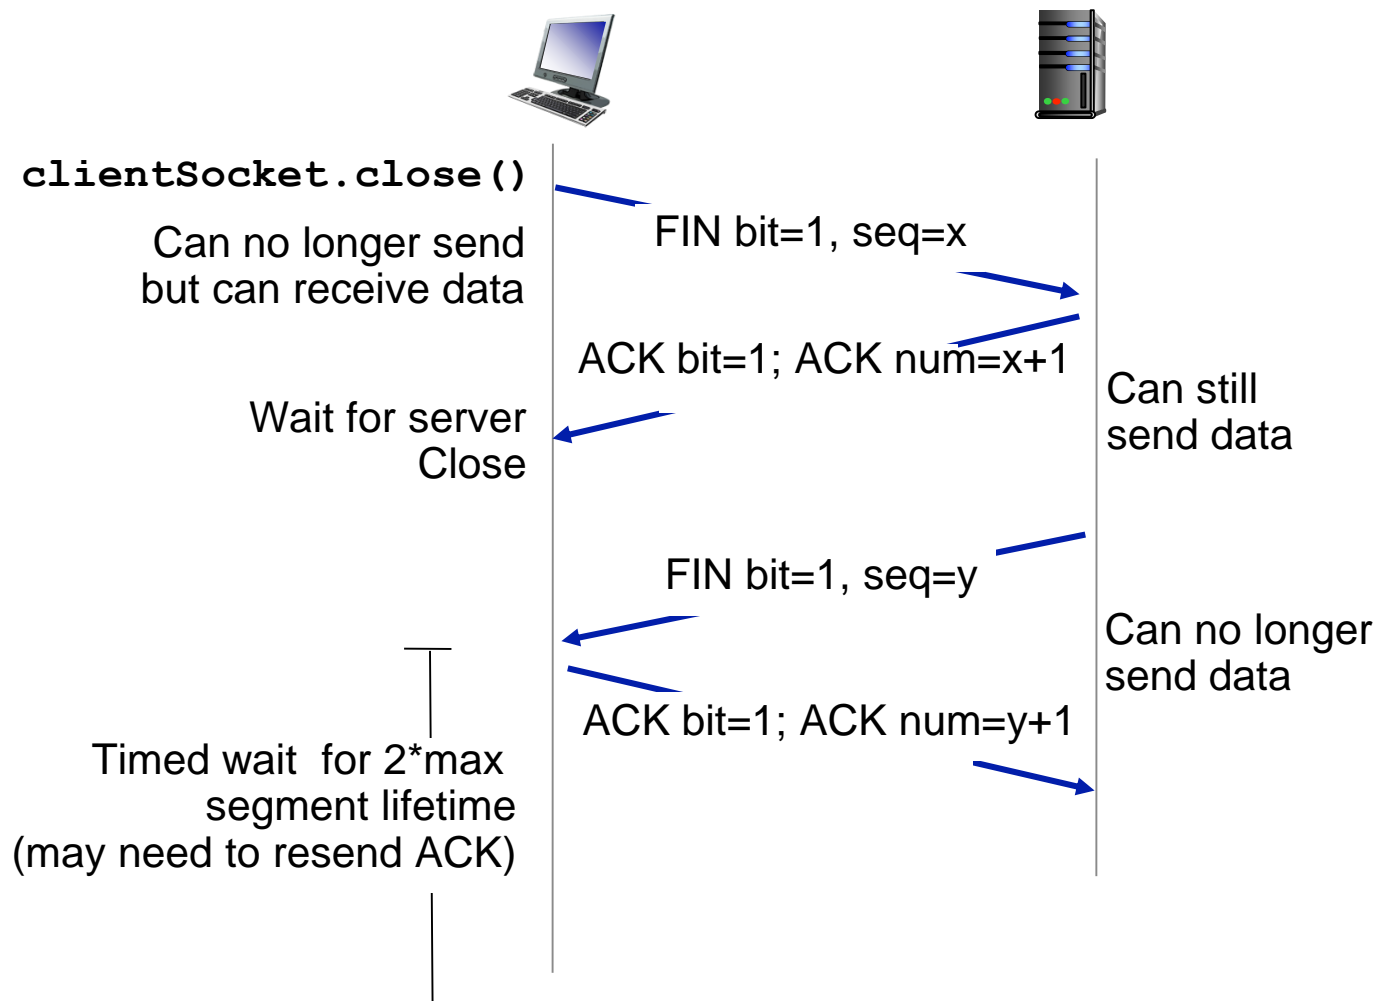
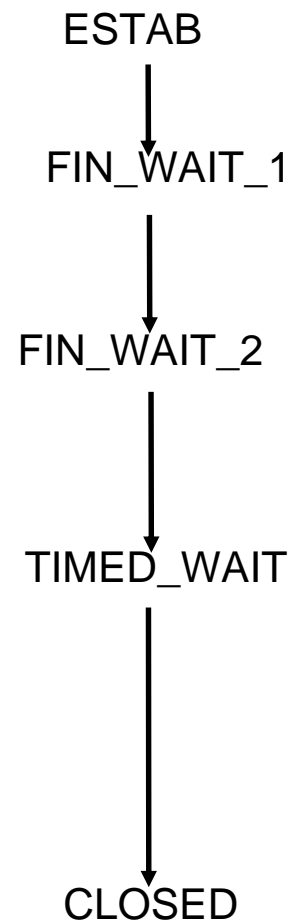
LISTEN
↓
SYN RCVD
↓
ESTAB

TCP: Closing a Connection


- **Client and server each close their side of the connection**
 - Send TCP segment with FIN bit = 1
- **Respond to received FIN with ACK**
 - On receiving FIN, ACK can be combined with own FIN
- **Simultaneous FIN exchanges can be handled**

TCP: Closing a Connection (Cont.)

Client State



Server State



Overview of Transport Layer

- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
 - Segment Structure
 - Reliable Data Transfer
 - Flow Control
 - Connection Management
- **Principles of Congestion Control**
- **TCP Congestion Control**

Principles of Congestion Control

- **What is congestion?**

- Informally: “too many sources sending too much data too fast for the *network* to handle”
- Different from flow control
 - Flow control used to ensure buffers at receiver do not overflow
 - Flow control does nothing to prevent router buffers from overflowing


- **What problems can congestion cause?**

- Lost packets (buffer overflow at routers)
- Long delays (queueing in router buffers)


Causes/Costs of Congestion: Scenario #1

- **Consider the following scenario**


- Two senders, two receivers
- One router with infinite buffers
- Single output link capacity shared by senders with capacity R
- Assume no retransmission necessary



- **Senders cannot send at a rate higher than $R/2$ since they are sharing single link**
- **As senders max out the output link, the delay between source and destination increases**



maximum per-connection throughput is $R/2$



large delays as arrival rate, λ_{in} , approaches capacity



Causes/Costs of Congestion: Scenario #2

- **Modify scenario #1 such that**

- The router has **finite shared buffers**
- The sender may retransmit packets

- **Retransmission reduce the throughput**

- Packets can be dropped at router if buffers are full
- Sender may timeout prematurely due to delay in router; send multiple copies of same data



Two Approaches Towards Congestion Control

- **End-to-end congestion control**

- No explicit feedback from network
- Congestion inferred from end-system observed loss, delay
- Approach taken by TCP

- **Network-assisted congestion control**

- Routers provide feedback to end systems
- Single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
- Explicit rate for sender to send at

Overview of Transport Layer

- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
- **Connection-oriented Transport: TCP**
 - Segment Structure
 - Reliable Data Transfer
 - Flow Control
 - Connection Management
- **Principles of Congestion Control**
- **TCP Congestion Control**


TCP Congestion Control

- **Must have each sender limit the rate at which it sends traffic as a function of the network congestion**
 - Too much congestion? Send less data.
 - Not much congestion? Full speed ahead!
- **How does a TCP sender limit the rate at which it sends data?**
- **How does a TCP sender detect congestion between itself and the destination?**
- **How should the sender change the rate at which it sends data based on the network congestion?**


TCP Congestion Control: AIMD

- **Additive Increase Multiplicative Decrease (AIMD)**

- **Approach**: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - **Additive increase**: increase congestion window (**cwnd**) by 1 MSS (Maximum Segment Size) every RTT until loss detected
 - ACKs arriving at sender signal the sender to increase its window
 - **Multiplicative decrease**: cut congestion window in half after a packet loss occurs



TCP Congestion Control (Cont.)



- **TCP sending rate (assuming no limit on receiver's buffer (rwnd))**

- Send `cwnd` bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

- **Congestion window (cwnd) is a dynamic, function of perceived network congestion**
 - Sender is limited by both `cwnd` and `rwnd`
- **Amount of unacked data at sender may not exceed the minimum of `cwnd` and `rwnd`**


$$\text{LastByteSent} - \text{LastByteAcked} \leq \min\{\text{cwnd}, \text{rwnd}\}$$

TCP Congestion-Control Algorithm

- **TCPs congestion control algorithm contain three main components**
 - Slow-start
 - Congestion Avoidance
 - Fast Recovery

TCP Slow-Start

- **When connection begins, increase rate exponentially until first loss event:**
 - Initially, `cwnd` = 1 MSS
 - Double `cwnd` every RTT
 - Done by incrementing `cwnd` for every ACK received
 - Initial rate is slow but ramps up exponentially fast
- **When first loss occurs, store ($.5 * \text{cwnd}$) as `SSThresh` and restart slow-start**
- **When `cwnd` reaches `SSThresh`, switch from slow-start mode to congestion avoidance mode**



Loss During Slow-Start

- **If loss is indicated by a timeout**

- Store $(.5 * \text{cwnd})$ as **SSThresh** and restart slow-start
- Set **cwnd** set to 1 MSS
- **cwnd** then grows exponentially (as in slow start) to threshold value **SSThresh**, then grows linearly in congestion avoidance phase

- **If loss is indicated by 3 duplicate ACKs (only in TCP Reno)**

- Duplicate ACKs indicate network capable of delivering some segments, so don't drop **cwnd** all the way down to 1 MSS
- Store $(.5 * \text{cwnd})$ as **SSThresh**
- **cwnd** is also set to $(.5 * \text{cwnd})$, but will be increment for each duplicate ACK

- **TCP Tahoe always sets cwnd to 1 (for either timeout or 3 duplicate acks)**

Switch from Slow Start to Congestion Avoidance

- Exponential growth phase shows TCP slow-start
- Linear phase after crossing over `ssThresh` shows the congestion avoidance phase
- TCP Tahoe
 - Set `cwnd` = 1 for both a timeout and for triple duplicate ACKs
 - Set `ssThresh` = `cwnd`/2
 - Re-enters slow-start phase
- TCP Reno
 - Implements Fast Recovery
 - Retransmits missing segment
 - Set `ssThresh` = `cwnd`/2
 - Set `cwnd` = `ssThresh` + 3
 - In congestion avoidance phase

